

Destination calculus

A linear λ -calculus for pure, functional memory updates

ANONYMOUS AUTHOR(S)

Destination-passing—aka out-parameters—is taking a parameter to fill rather than returning a result from a function. Due to its apparent imperative nature, destination passing has struggled to find its way to pure functional programming. In this paper, we present a pure core calculus with destinations. Our calculus subsumes all the existing systems, and can be used to reason about their correctness or extension. In addition our calculus can express programs that were previously not known to be expressible in a pure language. This is guaranteed by a modal type system where modes are used to represent both linear types and a system of ages to manage scopes. Type safety of our core calculus was largely proved formally with the Coq proof assistant.

CCS Concepts: • **Theory of computation** → **Type structures**; • **Software and its engineering** → **Formal language definitions**; **Functional languages**; **Data types and structures**.

Additional Key Words and Phrases: destination, functional programming, linear types, pure language

ACM Reference Format:

Anonymous Author(s). 2025. Destination calculus: A linear λ -calculus for pure, functional memory updates. *Proc. ACM Program. Lang.* XX, XX, Article XXX (January 2025), 28 pages. <https://doi.org/XXXXXXX.XXXXXXX>

1 INTRODUCTION

In destination-passing style a function doesn't return a value: it takes as an argument a location where the value ought to be returned. In our notation, a function of type $T \rightarrow U$ would, in destination-passing style, have type $T \rightarrow [U] \rightarrow 1$ instead. This style is common in systems programming, where destinations $[U]$ are more commonly known as “out parameters”. In C, $[U]$ would typically be a pointer of type U^* .

The reason why system programs rely on destinations so much is that using destinations can save calls to the memory allocator. If a function returns a U , it has to allocate the space for a U . But with destinations, the caller is responsible for finding space for a U . The caller may simply ask for the space to the memory allocator, in which case we've saved nothing; but it can also reuse the space of an existing U which it doesn't need anymore, or it could use a space in an array, or it could allocate the space in a region of memory that the memory allocator doesn't have access to, like a memory-mapped file.

This does all sound quite imperative, but we argue that the same considerations are relevant for functional programming, albeit to a lesser extent. In fact [Shaikhha et al. 2017] has demonstrated that using destination passing in the intermediate language of a functional array-programming language allowed for some significant optimizations. Where destinations truly shine in functional programming, however, is that they increase the expressiveness of the language; destinations as first-class values allow for meaningfully new programs to be written. This point was first explored in [Bagrel 2024].

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2025 Copyright held by the owner/author(s). Publication rights licensed to ACM.

2475-1421/2025/1-ARTXXX \$15.00

<https://doi.org/XXXXXXX.XXXXXXX>

The trouble, of course, is that destinations are imperative; we wouldn't want to sacrifice the immutability of our linked data structure (we'll usually just say "structure") for the sake of the more situational destinations. The goal is to extend functional programming just enough to be able to build immutable structures by destination passing without endangering purity and memory safety. This is precisely what [Bagrel 2024] does, using a linear type system to restrict mutation. Destinations become write-once references into an immutable structure with holes. In that we follow their leads, but we refine the type system further to allow for even more programs, as we discuss in Section 3.

There are two key elements to the expressiveness of destination- passing:

- structures can be built in any order. Not only from the leaves to the root, like in ordinary functional programming, but also from the root to the leaves, or any combination thereof. This can be done in ordinary functional program using function composition in a form of continuation-passing, destinations act as an optimization. This line of work was pioneered by [Minamide 1998]. While this only increases expressiveness when combined with the next point, the optimization is significant enough that destination passing has been implemented in the Ocaml optimizer to support tail modulo constructor [Bour et al. 2021];
- destinations are first-class values, which can be passed and stored like ordinary values. This is the innovation of [Bagrel 2024] upon which we build. The consequence is that not only the order in which a structure is built is arbitrary, this order can be determined dynamically during the runtime of the program.

To support this programming style, we introduce λ_d . We intend λ_d to serve as a core calculus to reason about safe destinations. Indeed λ_d subsumes all the systems that we've discussed in this section: they can all be encoded in λ_d via simple macro expansion. As such we expect that potential extensions to these systems can be justified by giving their semantics as an expansion in λ_d .

Our contributions are as follows:

- λ_d , a linear and modal simply typed λ -calculus with destinations (Sections 5 and 6). λ_d is expressive enough so that previous calculi for destinations can be encoded in λ_d (see Section 9);
- a demonstration that λ_d is more expressive than previous calculi with destinations (Sections 3 and 4), namely that destinations can be stored in structures with holes. We show how we can improve, in particular, on the breadth-first traversal example of [Bagrel 2024];
- an implementation strategy for λ_d which uses mutation without compromising the purity of λ_d (Section 8);
- formally-verified proofs, with the Coq proof assistant, of the main safety lemmas (Section 7).

2 WORKING WITH DESTINATIONS

Let's introduce and get familiar with λ_d , our simply typed λ -calculus with destination. The syntax is standard, except that we use linear logic's $\mathsf{T} \oplus \mathsf{U}$ and $\mathsf{T} \otimes \mathsf{U}$ for sums and products, since λ_d is linearly typed, even though it isn't a focus in this section.

2.1 Building up a vocabulary

In its simplest form, destination passing, much like continuation passing, is using a location, received as an argument, to return a value. Instead of a function with signature $\mathsf{T} \rightarrow \mathsf{U}$, in λ_d you would have $\mathsf{T} \rightarrow [\mathsf{U}] \rightarrow 1$, where $[\mathsf{U}]$ is read "destination for type U ". For instance, here is a destination-passing version of the identity function:

$$\begin{aligned} \text{dId} &: \mathsf{T} \rightarrow [\mathsf{T}] \rightarrow 1 \\ \text{dId } x \ d &\triangleq d \blacktriangleleft x \end{aligned}$$

We think of a destination as a reference to an uninitialized memory location, and $d \blacktriangleleft x$ (read “fill d with x ”) as writing x to the memory location.

The form $d \blacktriangleleft x$ is the simplest way to use a destination. But we don’t have to fill a destination with a complete value in a single step. Destinations can be filled piecemeal.

fillWithInlCtor : $[T \oplus U] \rightarrow [T]$
fillWithInlCtor $d \triangleq d \blacktriangleleft \text{Inl}$

In this example, we’re filling a destination for type $T \oplus U$ by setting the outermost constructor to left variant **Inl**. We think of $d \blacktriangleleft \text{Inl}$ (read “fill d with **Inl**”) as allocating memory to store a block of the form **Inl** \square , write the address of that block to the location that d points to, and return a new destination of type $[T]$ pointing to the uninitialized argument of **Inl**. Uninitialized memory, when part of a structure or value, like \square in **Inl** \square , is called a *hole*.

Notice that with **fillWithInlCtor** we are constructing the structure from the outermost constructor inward: we’ve written a value of the form **Inl** \square into a hole, but we have yet to describe what goes in the new hole \square . Such data constructors with uninitialized arguments are called *hollow constructors*. This is opposite to how functional programming usually works, where values are built from the innermost constructors outward: first we make a value v and only then can we use **Inl** to make an **Inl** v . This will turn out to be a key ingredient in the expressiveness of destination passing.

Yet, everything we’ve shown so far could have been done with continuations. So it’s worth asking: how are destination different from continuations? Part of the answer lies in our intention to effectively implement destinations as pointers to uninitialized memory (see Section 8). But where destinations really differ from continuations is when one has several destinations at hand. Then they have to fill *all* the destinations; whereas when one has multiple continuations, they can only return to one of them. Multiple destination arises when a destination of pair gets filled with a hollow pair constructor:

fillWithPairCtor : $[T \otimes U] \rightarrow [T] \otimes [U]$
fillWithPairCtor $d \triangleq d \blacktriangleleft (,)$

After using **fillWithPairCtor**, the user must fill both the first field *and* the second field, using the destinations of type $[T]$ and $[U]$ respectively. In plain English, it sounds obvious, but the key remark is that **fillWithPairCtor** doesn’t exist on continuations.

Structures with holes. It is crucial to note that while a destination is used to build a structure, the type of the structure being built might be different from the type of the destination that is being filled. A destination of type $[T]$ is a pointer to a yet-undefined part of a bigger structure. We say that such a structure has a hole of type T ; but the type of the structure itself isn’t specified (and never appears in the signature of destination-filling functions). For instance, using **fillWithPairCtor** only indicates that the structure being operated on has a hole of type $T \otimes U$ that is being written to.

Thus, we still need a type to tie the structure under construction — left implicit by destination-filling primitives — with the destinations representing its holes. To represent this, λ_d introduces a type $S \ltimes [T]$ for a structure of type S missing a value of type T to be complete. There can be several holes in S , resulting in several destinations on the right hand side: for example, $S \ltimes ([T] \otimes [U])$ represents a S that misses both a T and a U to be complete.

The general form $S \ltimes T$ is read “ S ampar T ”. The name “ampar” stands for “asymmetric memory par”; we will explain why it is asymmetric in Section 5.2. For now, it’s sufficient to observe that $S \ltimes [T]$ is akin to a “par” type $S \wp T^\perp$ in linear logic; you can think of $S \ltimes [T]$ as a (linear) function from T to S . That structures with holes could be seen as linear functions was first observed in [Minamide 1998], we elaborate on the value of having a par type rather than a function type in Section 4. A similar connective is called **Incomplete** in [Bagrel 2024].

Destinations always exist within the context of a structure with holes. A destination is both a witness of a hole present in the structure, and a handle to write to it. Crucially, destinations are otherwise ordinary values. To access the destinations of an ampar, λ_d provides a **map** construction, which lets us apply a function to the right-hand side of the ampar. It is in the body of the **map** construction that functions operating on destinations can be called:

```

fillWithInlCtor' : S × [T ⊕ U] → S × [T]
fillWithInlCtor' x ≜ map x with d ↦ fillWithInlCtor d
fillWithPairCtor' : S × [T ⊗ U] → S × ([T] ⊗ [U])
fillWithPairCtor' x ≜ map x with d ↦ fillWithPairCtor d

```

To tie this up, we need a way to introduce and to eliminate structures with holes. Structures with holes are introduced with **alloc** which creates a value of type $T \times [T]$. **alloc** is a bit like the identity function: it is a hole (of type T) that needs a value of type T to be a complete value of type T . Memory-wise, it is an uninitialized block large enough to host a value of type T , and a destination pointing to it. Conversely, structures with holes are eliminated with¹ **from_K** : $S \times 1 \rightarrow S$: if all the destinations have been consumed and only unit remains on the right side, then S no longer has holes and thus is just a normal, complete structure.

Equipped with these, we can, for instance, derive traditional constructors from piecemeal filling. In fact, λ_d doesn't have primitive constructor forms, constructors in λ_d are syntactic sugar. We show here the definition of **Inl** and **(,)**, but the other constructors are derived similarly.

```

Inl : T → T ⊕ U
Inl x ≜ from'_K (map alloc with d ↦ d ◁ Inl ◁ x)

(,) : T → U → T ⊗ U
(x, y) ≜ from'_K (map alloc with d ↦ case (d ◁ (,)) of (d1, d2) ↦ d1 ◁ x ; d2 ◁ y)

```

Memory safety and purity. At this point, the reader may be forgiven for feeling distressed at all the talk of mutations and uninitialized memory. How is it consistent with our claim to be building a pure and memory-safe language? The answer is that it wouldn't be if we'd allow unrestricted use of destination. Instead λ_d uses a linear type system to ensure that:

- destination are written at least once, preventing examples like:

```

forget : T
forget ≜ from'_K (map alloc with d ↦ ())

```

where reading the result of **forget** would result in reading the location pointed to by a destination that we never used, in other words, reading uninitialized memory;

- destination are written at most once, preventing examples like:

```

ambiguous1 : Bool
ambiguous1 ≜ from'_K (map alloc with d ↦ d ◁ true ; d ◁ false)

ambiguous2 : Bool
ambiguous2 ≜ from'_K (map alloc with d ↦ let x := (d ◁ false) in d ◁ true ; x)

```

where **ambiguous1** returns false and **ambiguous2** returns true due to evaluation order, even though let-expansion should be valid in a pure language.

¹As the name suggest, there is a more general elimination **from_K**. It will be discussed in Section 5.

<pre> 197 DList T \triangleq (List T) \ltimes [List T] 198 append : DList T \rightarrow T \rightarrow DList T 199 ys append y \triangleq 200 map ys with dys \mapsto case (dys \triangleleft (::)) of 201 (dy, dys') \mapsto dy \blacktriangleleft y \circ dys' 202 concat : DList T \rightarrow DList T \rightarrow DList T 203 ys concat ys' \triangleq map ys with d \mapsto d \circ ys' 204 toList : DList T \rightarrow List T 205 toList ys \triangleq from'_{\ltimes} (map ys with d \mapsto d \triangleleft []) </pre>	<pre> Queue T \triangleq (List T) \otimes (DList T) singleton : T \rightarrow Queue T singleton x \triangleq (Inr (x :: []), alloc) enqueue : Queue T \rightarrow T \rightarrow Queue T q enqueue y \triangleq case q of (xs, ys) \mapsto (xs, ys append y) dequeue : Queue T \rightarrow 1 \oplus (T \otimes (Queue T)) dequeue q \triangleq case q of { ((x :: xs), ys) \mapsto Inr (x, (xs, ys)), ([], ys) \mapsto case (toList ys) of { [] \mapsto Inl (), x :: xs \mapsto Inr (x, (xs, alloc)) } } </pre>
---	--

Fig. 1. Difference list and queue implementation in equirecursive λ_d

2.2 Functional queues, with destinations

Now that we have an intuition of how destinations work, let's see how they can be used to build usual data structures. For that section, we suppose that λ_d is equipped with equirecursive types and a fixed-point operator, that isn't part of our formally proven fragment.

Linked lists. We define lists as the fixpoint of the functor $X \mapsto 1 \oplus (T \otimes X)$. For convenience, we also define synthetic filling operators $\triangleleft []$ and $\triangleleft (::)$:

<pre> 221 List T $\triangleq^{\text{rec}} 1 \oplus (T \otimes (\text{List T}))$ 222 $\triangleleft [] : [\text{List T}] \rightarrow 1$ 223 d $\triangleleft [] \triangleq$ d \triangleleft Inl \triangleleft () </pre>	<pre> 224 $\triangleleft (::) : [\text{List T}] \rightarrow [\text{T}] \otimes [\text{List T}]$ 225 d $\triangleleft (::) \triangleq$ d \triangleleft Inr \triangleleft () </pre>
---	---

Just like we did in Section 2.1 we can recover traditional constructors from filling operators:

```

226   (::) : T  $\otimes$  (List T)  $\rightarrow$  List T
227   x :: xs  $\triangleq$  from'_{\ltimes} (map alloc with d  $\mapsto$  case (d  $\triangleleft$  (::)) of (dx, dxs)  $\mapsto$  dx  $\blacktriangleleft$  x  $\circ$  dxs  $\triangleleft$  xs)

```

Difference lists. Just like in any language, iterated concatenation of lists $((xs_1 ++ xs_2) ++ \dots) ++ xs_n$ is quadratic in λ_d . The usual solution to this is difference lists. The name difference lists covers many related implementation, but in pure functional languages, a difference list is usually represented as a function [Hughes 1986]. A singleton difference list is $\lambda ys \mapsto x :: ys$, and concatenation of difference lists is function composition. Difference lists are turned into a list by applying it to the empty list. The consequence is that no matter how many composition we have, each cons cell will be allocated a single time, making the iterated concatenation linear indeed.

However, each concatenation allocates a closure. If we're building a difference list from singletons and composition, there's roughly one composition per cons cell, so iterated composition effectively performs two traversals of the list. We can do better!

In λ_d we can represent a difference list as a list with a hole. A singleton difference list is $x :: \square$, concatenation is filling the hole with another difference list. The details are on the left of Figure 1. This encoding makes no superfluous traversal; in fact, concatenation is an $O(1)$ in-place update.

Efficient queue using previously defined structures. A simple implementation of a queue in a purely functional language is as a pair of lists (*front*, *back*) [Hood and Melville 1981]. Such queues are called *Hood-Melville queues*. Elements are popped from *front* and are enqueued in *back*. When we need to pop an element and *front* is empty, then we set the queue to (**reverse** *back*, []), and pop from the new front.

For such a simple implementation, Hood-Melville queues are surprisingly efficient: the cost of the reverse operation is $O(1)$ amortized for a single-threaded use of the queue. Still, it would be better to get rid of this full traversal of the back list.

Taking a step back, this *back* list that has to be reversed before it is accessed is really merely a representation of a list that can be extended from the back. And we already know an efficient implementation of lists that can be extended from the back but only accessed in a single-threaded fashion: difference lists.

So we can give an improved version of the simple functional queue using destinations. This implementation is presented on the right-hand side of Figure 1. Note that contrary to an imperative programming language, we can't implement the queue as a single difference list: our type system prevents us from reading the front elements of difference lists. Just like for the simple functional queue, we need a pair of one list that we can read from, and one that we can extend. Nevertheless this implementation of queues is both pure, as guaranteed by the λ_d type system, and nearly as efficient as what an imperative programming language would afford.

3 SCOPE ESCAPE OF DESTINATIONS

In Section 2, we've silently been making an assumption: establishing a linear discipline on destinations ensures that all destinations will eventually find their way to the left of a fill operator \blacktriangleleft or \triangleleft , so that the associated holes get written to. This turns out to be more subtle than it may first appear.

To see why, let's consider the type $[[T]]$: the type of a destination pointing to a hole where a destination is expected. Think of it as an equivalent of the pointer type $T **$ in the C language. Destinations are indeed ordinary values, so they can be stored in data structures, and before they get effectively stored, holes stands in their places in the structure. In particular if we have $d : [T]$ and $dd : [[T]]$, we can form $dd \blacktriangleleft d$.

This, by itself, is fine: we're building a linear data structure containing a destination; the destination d is used linearly as it will eventually be consumed when that structure is consumed. However, as we explained in Section 2.1, destinations always exist within the scope of a structure with holes, and they witness how incomplete the structure is. As a result, to make a structure readable, it is not enough to know that all its remaining destinations will end up on the left of a fill operator *eventually*; they must do so *now*, before the scope they belong to ends. Otherwise some holes might not have been written to *yet* when the structure is made readable.

The problem is, with a malicious use of $dd \blacktriangleleft d$, we can store away d in a parent scope, so d might not end up on the left of a fill operator before the scope it originates from is complete. And still, that would be accepted by the linear type system.

To visualize the problem, let's consider simple linear store semantics. We'll need the **alloc'** operator²

$$\begin{aligned} \text{alloc}' &: ([T] \rightarrow 1) \rightarrow T \\ \text{alloc}' f &\triangleq \text{from}'_{\mathbf{x}} (\text{map } \text{alloc } \text{with } d \mapsto f d) \end{aligned}$$

²The actual semantics that we'll develop in Section 6 is more refined and can give a semantics to the more flexible $\text{from}'_{\mathbf{x}} t$ and **alloc** operators directly, but for now, this simpler semantic will suffice for **alloc'**.

The semantics of **alloc'** is: allocate a hole in the store, call the function with the corresponding destination; when the function has returned, dereference the destination to obtain a **T**. Which we can visualize as:

$$\mathcal{S} \mid \mathbf{alloc}' (\lambda d \mapsto t) \longrightarrow \mathcal{S} \sqcup \{h := \square\} \mid (t[d := \rightarrow h] \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow h)$$

For instance:

$$\begin{aligned} & \{\} \mid \mathbf{alloc}' (\lambda d \mapsto d \triangleleft \text{Inl} \triangleleft ()) \\ \longrightarrow & \{h := \square\} \mid \rightarrow h \triangleleft \text{Inl} \triangleleft () \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow h \\ \longrightarrow & \{h := \text{Inl} ()\} \mid \mathbf{deref} \rightarrow h \\ \longrightarrow & \{\} \mid \text{Inl} () \end{aligned}$$

Now, we are ready to see a counterexample and how it goes wrong:

$$\mathbf{alloc}' (\lambda dd \mapsto \mathbf{case} (\mathbf{alloc}' (\lambda d \mapsto dd \triangleleft d)) \mathbf{of} \{ \text{true} \mapsto (), \text{false} \mapsto () \})$$

Here $dd : \llbracket \text{Bool} \rrbracket$ and $d : \llbracket \text{Bool} \rrbracket$. The problem with this example stems from the fact that d is fed to dd , instead of being filled, in the scope of d . Let's look at the problem in action:

$$\begin{aligned} & \{\} \mid \mathbf{alloc}' (\lambda dd \mapsto \mathbf{case} (\mathbf{alloc}' (\lambda d \mapsto dd \triangleleft d)) \mathbf{of} \{ \text{true} \mapsto (), \text{false} \mapsto () \}) \\ \longrightarrow & \{hd := \square\} \mid \mathbf{case} (\mathbf{alloc}' (\lambda d \mapsto \rightarrow hd \triangleleft d)) \mathbf{of} \{ \text{true} \mapsto (), \text{false} \mapsto () \} \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow hd \\ \longrightarrow & \{hd := \square, h := \square\} \mid \mathbf{case} (\rightarrow hd \triangleleft \rightarrow h \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow h) \mathbf{of} \{ \text{true} \mapsto (), \text{false} \mapsto () \} \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow hd \\ \longrightarrow & \{hd := \rightarrow h, h := \square\} \mid \mathbf{case} (\mathbf{deref} \rightarrow h) \mathbf{of} \{ \text{true} \mapsto (), \text{false} \mapsto () \} \mathbin{\text{\texttt{;}}} \mathbf{deref} \rightarrow hd \end{aligned}$$

Because of d escaping its scope, we end up reading uninitialized memory.

This example must be rejected by our type system. As demonstrated by [Bagrel 2024], it's possible to reject this example using purely a linear type system: they make it so that, in $d \triangleleft t$, t can't be linear. Since all destinations are linear, t can't be, or contain, a destination. This is a rather blunt restriction! Indeed, it makes the type $\llbracket \text{T} \rrbracket$ practically useless: we can form destinations with type $\llbracket \text{T} \rrbracket$ but they can never be filled. More concretely, this means that destination can never be stored in data structures with holes, such as the difference list or queues of Section 2.2.

But we really want to be able to store destinations in data structures with holes, in fact, we'll use this capability to our advantage in Figure 2. So we want t in $d \triangleleft t$ to be allowed to be linear. Without further restriction, the counterexample would be well-typed. To address this λ_d uses a system of ages to represent scope. Ages are described in Section 5.

4 BREADTH-FIRST TREE TRAVERSAL

As a more full-fledged example, which uses the full expressive power of λ_d , we borrow and improve on an example from [Bagrel 2024], breadth-first tree relabeling:

Given a tree, create a new one of the same shape, but with the values at the nodes replaced by the numbers $1 \dots |T|$ in breadth-first order.

This example cannot be implemented using [Minamide 1998] system where structures with holes are represented as linear functions. Destinations as first-class values are very much required. Indeed, breadth-first traversal implies that the order in which the structure must be populated (left-to-right, top-to-bottom) is not the same as the structural order of a functional binary tree, that is, building the leaves first and going up to the root. This isn't very natural in functional programming, which is why a lot has been written on purely functional breadth-first traversals [Gibbons 1993; Gibbons et al. 2023; Okasaki 2000].

On the other hand, as demonstrated in [Bagrel 2024], first-class destination passing lets us use the familiar algorithm with queues that is taught in classroom. Specifically, the algorithm keeps a queue of pairs (*input subtree*, *destination to output subtree*).

Figure 2 presents the λ_d implementation of the breadth-first tree traversal. There, **Tree T** is defined unsurprisingly as $\mathbf{Tree\ T} \triangleq 1 \oplus (\mathbf{T} \otimes ((\mathbf{Tree\ T}) \otimes (\mathbf{Tree\ T})))$; we refer to the synthetic constructors of

There's an error left in the implementation: the state is specified to be linear, but the natural number state is used non-linearly. Fix

```

344 go : ( $S_{100} \rightarrow T_1 \rightarrow (!_{100} S) \otimes T_2$ )  $\omega_{100} \rightarrow S_{100} \rightarrow \text{Queue } (\text{Tree } T_1 \otimes \text{Tree } T_2) \rightarrow (!_{100} S)$ 
345 go  $f$   $st$   $q \triangleq_{\text{rec}}$  case (dequeue  $q$ ) of {
346   Inl ()  $\mapsto E_{100} st$ ,
347   Inr (( $tree, dtree$ ),  $q'$ )  $\mapsto$  case  $tree$  of {
348     Nil  $\mapsto dtree \triangleleft \text{Nil} \ ; \ \text{go } f$   $st$   $q'$ ,
349     Node  $x$   $tl$   $tr \mapsto$  case ( $dtree \triangleleft \text{Node}$ ) of
350       ( $dy, (dtr, dtr)$ )  $\mapsto$  case ( $f$   $st$   $x$ ) of
351         ( $E_{100} st', y$ )  $\mapsto$ 
352            $dy \blacktriangleleft y \ ;$ 
353           go  $f$   $st'$  ( $q'$  enqueue ( $tl, dtr$ ) enqueue ( $tr, dtr$ ))
354       }
355   }
356 mapAccumBFS : ( $S_{100} \rightarrow T_1 \rightarrow (!_{100} S) \otimes T_2$ )  $\omega_{100} \rightarrow S_{100} \rightarrow \text{Tree } T_1 \rightarrow \text{Tree } T_2 \otimes (!_{100} S)$ 
357 mapAccumBFS  $f$   $st$   $tree \triangleq$  from $\kappa$  (map alloc with  $dtree \mapsto$  go  $f$   $st$  (singleton ( $tree, dtree$ )))
358 relabelDPS :  $\text{Tree } 1_{100} \rightarrow (\text{Tree } \text{Nat}) \otimes (!_{100} (!_{\omega V} \text{Nat}))$ 
359 relabelDPS  $tree \triangleq$  mapAccumBFS
360   ( $\lambda ex_{100} \mapsto \lambda un \mapsto un \ ; \ \text{case}_{100} ex$  of
361      $E_{\omega V} st \mapsto (E_{100} (E_{\omega V} (\text{succ } st)), st)$ )
362   ( $E_{\omega V} 1$ )
363    $tree$ 

```

Fig. 2. Breadth-first tree traversal in destination-passing style

Tree T as **Nil** and **Node**. We also assume some encoding of the type **Nat** of natural number. Queues are the efficient queues from Section 2.2.

New in Figure 2 are the **fuchsia** annotations: these are *modes*. We'll describe modes in detail in Section 5. In the meantime, **1** and ω control linearity: we use ω to mean that the function f is used many times.

TODO: Revisit starting from there, make state unrestricted and drop it at the end

; v and ∞ control age, we use ∞ to mean that neither the state nor the function f contain destinations, therefore they aren't restricted to a particular scope. Arguments with no modes are assumed to have mode **1v**: they're linear and can't escape the current scope. We also have the modality **!_mT**, which reifies mode **m** in a type, like exponential connectives in linear logic, **!_mT** is, by definition, such that $(!_{\text{m}} T) \rightarrow U \simeq T_{\text{m}} \rightarrow U$.

Figure 2 implements breadth-first relabeling **relabelDPS** as an instance of a more general breadth-first traversal function **mapAccumBFS**, which applies any state-passing style transformation of labels in breadth-first order, at the end of the processing.

In **mapAccumBFS**, we create a new destination $dtree$ into which we will write the result of the traversal, then call the main loop **go**. The **go** function is in destination-passing style, but what's remarkable is that **go** takes an unbounded number of destinations as arguments, since there are as many destinations as items in the queue. This is where we actually use the fact that destinations are ordinary values.

A similar breadth-first traversal function can be found in [Bagrel 2024]. The difference is that, because they can't store destinations in structures with holes (see the discussion in Section 3), their implementation can't use the efficient **Queue** T type from Section 2.2. So they have to revert to using a Hood-Melville queue for breadth-first traversal.

I'm not sure "current scope" is super clear, to be honest

we say in a footnote that **from** _{κ} will be introduced only in section 5, while we use it here without any explanations. - Arnaud: maybe let's use **from**' and drop the returned state from **mapAccumBFS** and **relabelDPS**? This removes the linearity of state, but maybe it's more legible?


```

393  t, u ::= x | t' t | t ; t'
394         | casem t of {Inl x1 ↦ u1, Inr x2 ↦ u2} | casem t of (x1, x2) ↦ u | casem t of En x ↦ u
395         | map t with x ↦ t' | toκ t | fromκ t | alloc
396         | t ◁ () | t ◁ Inl | t ◁ Inr | t ◁ (,) | t ◁ Em | t ◁ (λxm ↦ u) | t ◁ t'
397
398  v ::= h (hole)
399         | →h (destination)
400         | H⟨v2 ∧ v1⟩ (ampar value form)
401         | () | λxm ↦ u | Inl v | Inr v | Em v | (v1, v2)
402
403  T, U, S ::= [n T] (destination)
404              | S × T (ampar)
405              | 1 | T1 ⊕ T2 | T1 ⊗ T2 | !m T | Tm → U
406
407  m, n ::= pa (pair of multiplicity and age)
408  p ::= 1 | ω
409  a ::= v | ↑ | ∞
410
411  Ω, Γ, Θ, Δ ::= • | x :m T | h :n T | →h :m [n T]
412              | Ω1, Ω2 | Ω1 + Ω2 | m · Ω | →i Δ

```

Fig. 3. Grammar of λ_d

<pre> 415 t ◁ t' ≜ t ◁ (to_κ t') 416 Inl t ≜ from'_{κ} (map alloc with d ↦ 417 d ◁ Inl ◁ t 418) 419 Inr t ≜ from'_{κ} (map alloc with d ↦ 420 d ◁ Inr ◁ t 421) 422 E_m t ≜ from'_{κ} (map alloc with d ↦ 423 d ◁ E_m ◁ t 424) </pre>	<pre> 415 from'_{κ} t ≜ 416 case (from_{κ} (map t with un ↦ un ; E_{1∞} ())) of 417 (st, ex) ↦ case ex of 418 E_{1∞} un ↦ un ; st 419 420 λx_m ↦ u ≜ from'_{κ} (map alloc with d ↦ 421 d ◁ (λx_m ↦ u) 422) 423 424 (t₁, t₂) ≜ from'_{κ} (map alloc with d ↦ 425 case (d ◁ (,)) of 426 (d₁, d₂) ↦ d₁ ◁ t₁ ; d₂ ◁ t₂ 427) </pre>
---	---

Fig. 4. Syntactic sugar forms for terms

5 LANGUAGE SYNTAX AND TYPE SYSTEM

TODO: Move val grammar away?

TODO: We are missing syntactic sugar for some forms like unit () (needs DisposableOnly, do not forget)

λ_d is based on simply typed lambda calculus, with a first-order type system, featuring modal function types and modal boxing, in addition to unit (1), product (\otimes) and sum (\oplus) types. It is also equipped with the destination type $[_mT]$ and ampar type $S \times T$ that have been previewed in Section 2 to represent DPS structure building. The core grammar of the language is presented in Figure 3. We also provide commonly used syntactic sugar forms for terms in Figure 4.

Modes in λ_d have two axes — multiplicity (i.e. linear/non-linear), and age control — and they

Remove the terminology "axes" for the modes. This terminology is specific to the Oxidizing Ocaml work and,

+	\uparrow^n	∞	\cdot	\uparrow^n	∞	+	1	ω	\cdot	1	ω
\uparrow^m	if $n = m$ then \uparrow^n else ∞	∞	\uparrow^m	\uparrow^{n+m}	∞	1	ω	ω	1	1	∞
∞	∞	∞	∞	∞	∞	ω	ω	ω	ω	∞	∞

We pose $\uparrow^0 = \nu$ and $\uparrow^n = \uparrow \cdot \uparrow^{n-1}$

Fig. 5. Operation tables for age and multiplicity semirings

take place on variable bindings in typing contexts Ω , and on function arrows, but are not part of the type itself.

We omit the mode annotation \mathfrak{m} on function arrows and destinations when the mode in question is the multiplicative neutral element 1ν of the mode semiring (in particular, a function arrow without annotation is linear by default). A function arrow with multiplicity 1 is equivalent to the linear arrow \multimap from [Girard 1995].

Let's now introduce the age mode axis, which is a novel feature of this calculus.

5.1 Age-control for bindings to prevent scope escape of destinations

The solution we chose to alleviate scope escape of destinations (detailed in Section 3) is to track the age of destinations (as De-Brujin-like scope indices), and to set age-control restriction on the typing rule of destination-filling primitives.

Age is represented by a commutative semiring, where ν indicates that a destination originates from the current scope, and \uparrow indicates that it originates from the scope just before. Finally, age ∞ is introduced for variables standing in place of a non-age-controlled value. In particular, destinations can never have age ∞ ; the main role of age ∞ is thus to act as a proof that no destination can be part of the value.

Semiring addition $+$ is used to find the age of a variable or destination that is used in two subterms of a program. Semiring multiplication \cdot corresponds to age composition, and is in fact an integer sum on scope indices. ∞ is absorbing for both addition and multiplication. Tables for the operations $+$ and \cdot are presented in Figure 5.

The age commutative semiring is then combined with the multiplicity commutative semiring from [Bernardy et al. 2018] to form a canonical product commutative semiring that forms the mode of each typing context binding in our final type system.

5.2 Design motivation behind the ampar and destination types

Minamide's work [Minamide 1998] is the earliest record we could find of a functional calculus integrating the idea of incomplete data structures (structures with holes) that exist as first class values and can be interacted with by the user.

In that paper, a structure with a hole is named *hole abstraction*. In the body of a hole abstraction, the bound *hole variable* should be used linearly (exactly once), and must only be used as a parameter of a data constructor (it cannot be pattern-matched on). A hole abstraction of type $(\mathsf{T}, \mathsf{S})\mathsf{hfun}$ is thus a weak form of linear lambda abstraction $\mathsf{T} \multimap \mathsf{S}$, which just moves a piece of data into a bigger data structure.

Now, in classical linear logic, we know we can transform linear implication $\mathsf{T} \multimap \mathsf{S}$ into $\mathsf{S} \wp \mathsf{T}^\perp$. Doing so for the type $(\mathsf{T}, \mathsf{S})\mathsf{hfun}$ gives $\mathsf{S} \wp [\mathsf{T}]$, where $[\cdot]$ is memory negation, and \wp is a memory *par* (it allows less interaction than the CLL *par*, because hfun is weaker than \multimap).

Transforming the hole abstraction from its original implication form to a *par* form let us consider the *destination* type $[\mathsf{T}]$ as a first class component of our calculus. We also get to see the hole

You say "commutative semiring" a lot. But maybe it'll be relaxing to ponder that in reality, our semirings aren't really semirings at all: they don't have a 0. Instead, just say "semiring" and be clear about your definition. The first occurrence should mention (possibly in a footnote) that semirings in this paper aren't assume to have a 0 but are assumed to be commutative.

(Yann) I think they call this ringoid in the Bernardy et al. paper on modalities

I removed the bit about extending ages to variables because we, in fact, have no age (or multiplicity)

abstraction aka. memory par as a pair-like structure, where the two sides might be coupled together in a way that prevent using both of them simultaneously.

From memory par $\widehat{\bowtie}$ to ampar \ltimes . In CLL, thanks to the cut rule, any of the sides S or T of a par $S \bowtie T$ can be eliminated, by interaction with the opposite type \cdot^\perp , to free up the other side. But in λ_d , we have two types of interaction to consider: interaction between T and $\lfloor T \rfloor$, and interaction between T and $T \rightarrow \cdot$. The structure that may contain holes, S , can safely interact with $\lfloor S \rfloor$ (merge it into a bigger structure with holes), but not with $T \rightarrow \cdot$, as it would let the user read an incomplete structure! On the other hand, a complete value of type $T = (\dots \lfloor T' \rfloor \dots)$ containing destinations (but no holes) can safely interact with a function $T \rightarrow 1$ (in particular, the function can pattern-match on the value of type T to access the destinations), but it is not always safe to fill it into a $\lfloor T \rfloor$ as that might allow scope escape of destination $\lfloor T' \rfloor$ as we've just seen in Section 3.

To recover sensible rules for the connective, we decided to make it asymmetric, hence ampar $(S \ltimes T)$ for *asymmetrical memory par*:

- the left side S can contain holes, and can be only be eliminated by interaction with $\lfloor S \rfloor$ using “fillComp” (\Leftarrow) to free up the right side T ;
- the right side T cannot contain holes (it might contain destinations), and can be eliminated by interaction with $T \rightarrow 1$ to free up the left side S . At term level, this is done using **from** _{\ltimes} and **map**.

5.3 Typing of terms and values

The typing rules for λ_d are highly inspired from [Abel and Bernardy 2020] and Linear Haskell [Bernardy et al. 2018], and are detailed in Figure 6. In particular, we use the same additive/multiplicative approach on contexts for linearity and age enforcement

Destinations and holes are two faces of the same coin, as seen in Section 2.1, and must always be in 1:1 correspondence. Thus, the new idea of our type system is to feature *hole bindings* $\lfloor h \rfloor :_n T$ and *destination bindings* $\rightarrow h :_m \lfloor n T \rfloor$ in addition to the variable bindings $x :_m T$ that usually populates typing contexts.

Such bindings mention two distinct classes of names: regular variable names x, y , and *hole names* h, h_1, h_2 which are identifiers for a memory cell that hasn't been written to yet. Hole names are represented by natural numbers under the hood, so they are equipped with addition $h+h'$ and can act both as relative offsets or absolute positions in memory. Typically, when a structure is effectively allocated, its hole (and destination) names are shifted by the maximum hole name encountered so far in the program (denoted $\max(hnames(C))$); this corresponds to finding the next unused memory cell in which to write new data.

The mode n of a hole binding $\lfloor h \rfloor :_n T$ (also present in the corresponding destination type $\lfloor n T \rfloor$) indicates the mode a value must have to be written to it (that is to say, the mode of bindings that the value depends on to type correctly). We see the mode of a hole coming into play when a hole is located behind an exponential constructor: we should only write a non-linear value to the hole $\lfloor h \rfloor$ in $E_{\text{cov}} \lfloor h \rfloor$. In particular, we should not store a destination into this hole, otherwise it could later be extracted and used in a non-linear fashion.

A destination binding $\rightarrow h :_m \lfloor n T \rfloor$ mentions two modes, m and n ; m is the mode of the binding (in particular, it informs on the age of the destination itself). Whereas n is part of the destination's type $\lfloor n T \rfloor$ and corresponds to the mode a value has to have to be written to the corresponding hole. In a well-typed closed program the mode m of a destination binding can never be of multiplicity ω or age ∞ in the typing tree; it is always linear and of finite age.

We also extend mode product to a pointwise action on typing contexts:

Should I mention that dests is a non-involutive negation? – Arnaud: You should mention that dests aren't involutive. But I'm, as always, very opposed to call this sort of things negations. They're dualisations.

Thomas: why did you add hfill here (between the two parts of the figure)?

I only know what you mean by “must always be in 1:1 correspondence” because I already know what it's supposed to mean. Rephrase. Also, maybe it isn't the best place for it.

There's no hood for them to be under. Something like “For simplicity we'll assume that hole names are natural number”. But must explain why it simplifies anything.

This parenthetic remark doesn't read very well, I think that

$\Theta \vdash t : T$	(Typing judgment for terms)		
$\frac{\text{TY-TERM-VAR} \quad \text{DisposableOnly } \Theta \quad \textcolor{blue}{1}_V <: \textcolor{blue}{m}}{\Theta, x : \textcolor{blue}{m}T \vdash x : T}$	$\text{TY-TERM-APP} \quad \frac{\Theta_1 \vdash t : T \quad \Theta_2 \vdash t' : \textcolor{blue}{T}_m \rightarrow U}{\textcolor{blue}{m}\Theta_1 + \Theta_2 \vdash t' t : U}$	$\text{TY-TERM-PATU} \quad \frac{\Theta_1 \vdash t : 1 \quad \Theta_2 \vdash u : U}{\Theta_1 + \Theta_2 \vdash t \textcolor{blue}{; } u : U}$	
$\text{TY-TERM-PATS} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{T}_1 \oplus \textcolor{blue}{T}_2 \quad \Theta_2, x_1 : \textcolor{blue}{m}T_1 \vdash u_1 : U \quad \Theta_2, x_2 : \textcolor{blue}{m}T_2 \vdash u_2 : U}{\textcolor{blue}{m}\Theta_1 + \Theta_2 \vdash \text{case}_{\textcolor{blue}{m}} t \text{ of } \{ \text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2 \} : U}$	$\text{TY-TERM-PATP} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{T}_1 \otimes \textcolor{blue}{T}_2 \quad \Theta_2, x_1 : \textcolor{blue}{m}T_1, x_2 : \textcolor{blue}{m}T_2 \vdash u : U}{\textcolor{blue}{m}\Theta_1 + \Theta_2 \vdash \text{case}_{\textcolor{blue}{m}} t \text{ of } (x_1, x_2) \mapsto u : U}$		
$\text{TY-TERM-PATE} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{!}_n T \quad \Theta_2, x : \textcolor{blue}{m}_n T \vdash u : U}{\textcolor{blue}{m}\Theta_1 + \Theta_2 \vdash \text{case}_{\textcolor{blue}{m}} t \text{ of } E_n x \mapsto u : U}$	$\text{TY-TERM-MAP} \quad \frac{\Theta_1 \vdash t : U \times T \quad \textcolor{blue}{!}\uparrow \Theta_2, x : \textcolor{blue}{!}_V T \vdash t' : T'}{\Theta_1 + \Theta_2 \vdash \text{map } t \text{ with } x \mapsto t' : U \times T'}$		
$\text{TY-TERM-TOA} \quad \frac{\Theta \vdash u : U}{\Theta \vdash \text{to}_{\textcolor{blue}{K}} u : U \times 1}$	$\text{TY-TERM-FROMA} \quad \frac{\Theta \vdash t : U \times (\textcolor{blue}{!}_{\textcolor{blue}{IOO}} T)}{\Theta \vdash \text{from}_{\textcolor{blue}{K}} t : U \otimes (\textcolor{blue}{!}_{\textcolor{blue}{IOO}} T)}$	$\text{TY-TERM-ALLOC} \quad \frac{\text{DisposableOnly } \Theta}{\Theta \vdash \text{alloc} : T \times [T]}$	$\text{TY-TERM-FILLU} \quad \frac{\Theta \vdash t : \textcolor{blue}{!}_n 1}{\Theta \vdash t \triangleleft () : 1}$
$\text{TY-TERM-FILLL} \quad \frac{\Theta \vdash t : \textcolor{blue}{!}_n T_1 \oplus T_2}{\Theta \vdash t \triangleleft \text{Inl} : \textcolor{blue}{!}_n T_1}$	$\text{TY-TERM-FILLR} \quad \frac{\Theta \vdash t : \textcolor{blue}{!}_n T_1 \oplus T_2}{\Theta \vdash t \triangleleft \text{Inr} : \textcolor{blue}{!}_n T_2}$	$\text{TY-TERM-FILLP} \quad \frac{\Theta \vdash t : \textcolor{blue}{!}_n T_1 \otimes T_2}{\Theta \vdash t \triangleleft () : \textcolor{blue}{!}_n T_1 \otimes \textcolor{blue}{!}_n T_2}$	$\text{TY-TERM-FILLE} \quad \frac{\Theta \vdash t : \textcolor{blue}{!}_n \textcolor{blue}{n}' T}{\Theta \vdash t \triangleleft E_{\textcolor{blue}{n}'} : \textcolor{blue}{!}_{\textcolor{blue}{n}' n} T}$
$\text{TY-TERM-FILLF} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{!}_n T_m \rightarrow U \quad \Theta_2, x : \textcolor{blue}{m}T \vdash u : U}{\Theta_1 + (\textcolor{blue}{!}\uparrow \cdot \textcolor{blue}{n}) \cdot \Theta_2 \vdash t \triangleleft (\lambda x \textcolor{blue}{m} \mapsto u) : 1}$	$\text{TY-TERM-FILLCOMP} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{!}_n U \quad \Theta_2 \vdash t' : U \times T}{\Theta_1 + (\textcolor{blue}{!}\uparrow \cdot \textcolor{blue}{n}) \cdot \Theta_2 \vdash t \triangleleft \textcolor{blue}{\circ} t' : T}$		
$\Theta \vdash t : T$	(Derived typing judgment for syntactic sugar forms)		
$\text{TY-STERM-FROMA}' \quad \frac{\Theta \vdash t : T \times 1}{\Theta \vdash \text{from}'_{\textcolor{blue}{K}} t : T}$	$\text{TY-STERM-FILLLEAF} \quad \frac{\Theta_1 \vdash t : \textcolor{blue}{!}_n T \quad \Theta_2 \vdash t' : T}{\Theta_1 + (\textcolor{blue}{!}\uparrow \cdot \textcolor{blue}{n}) \cdot \Theta_2 \vdash t \triangleleft t' : 1}$	$\text{TY-STERM-FUN} \quad \frac{\Theta_2, x : \textcolor{blue}{m}T \vdash u : U}{\Theta_2 \vdash \lambda x \textcolor{blue}{m} \mapsto u : T_m \rightarrow U}$	
$\text{TY-STERM-LEFT} \quad \frac{\Theta_2 \vdash t : T_1}{\Theta_2 \vdash \text{Inl } t : T_1 \oplus T_2}$	$\text{TY-STERM-RIGHT} \quad \frac{\Theta_2 \vdash t : T_2}{\Theta_2 \vdash \text{Inr } t : T_1 \oplus T_2}$	$\text{TY-STERM-EXP} \quad \frac{\Theta_2 \vdash t : T}{\textcolor{blue}{m}\Theta_2 \vdash E_{\textcolor{blue}{m}} t : \textcolor{blue}{!}_m T}$	$\text{TY-STERM-PROD} \quad \frac{\Theta_{21} \vdash t_1 : T_1 \quad \Theta_{22} \vdash t_2 : T_2}{\Theta_{21} + \Theta_{22} \vdash (t_1, t_2) : T_1 \otimes T_2}$

Fig. 6. Typing rules for terms and syntactic sugar

$$\begin{cases} \textcolor{blue}{n}' \cdot (x : \textcolor{blue}{m}T) &= x : \textcolor{blue}{n}' \cdot \textcolor{blue}{m}T \\ \textcolor{blue}{n}' \cdot (\textcolor{blue}{!}h : T) &= \textcolor{blue}{!}h : \textcolor{blue}{n}' \cdot n T \\ \textcolor{blue}{n}' \cdot (\rightarrow h : \textcolor{blue}{m} \textcolor{blue}{!}_n T) &= \rightarrow h : \textcolor{blue}{n}' \cdot \textcolor{blue}{m} \textcolor{blue}{!}_n T \end{cases}$$

Figure 6 presents the typing rules for terms, and rules for syntactic sugar forms that have been derived from term rules and proven formally too. Figure 8 presents the typing rules for values of the language. In every figure,

- Ω denotes an arbitrary typing context, with no particular constraint;

- Γ denotes a typing context made only of hole and destination bindings;
- Θ denotes a typing context made only of destination and variable bindings;
- Δ denotes a typing context made only of destination bindings.

5.3.1 Typing of terms \vdash . A term t always types in a context Θ made only of destination and variable bindings. That being said, typing rules for terms and their syntactic sugar in Figure 6 never explicitly mention a destination binding; only variable bindings. Only at runtime some variables will be substituted by destinations having a matching type and mode (that is why terms cannot type in a context made of variable bindings alone). As a result, the user has to deal with is only slightly more complex than a linear type system *à la* [Bernardy et al. 2018].

Let's focus on a few peculiarities of the type system for terms. The predicate `DisposableOnly` Θ in rules `TY-TERM-VAL` and `TY-TERM-VAR` says that Θ can only contain variable bindings with multiplicity ω , for which weakening is allowed in linear logic. It is enough to allow weakening at the leaves of the typing tree, that is to say the two aforementioned rules (`TY-TERM-VAL` is indeed a leaf for judgments \vdash , that holds a subtree of judgments $\nabla\vdash$).

Rule `TY-TERM-VAR`, in addition to weakening, allows for dereliction of the mode for the variable used, with subtyping constraint $1\nu <: m$ defined as $pa <: p'a' \iff p <: p' \wedge a <: a'$ where:

$$\left\{ \begin{array}{l} 1 <: 1 \\ p <: \omega \end{array} \right\} \quad \left\{ \begin{array}{l} \uparrow^m <: \uparrow^n \iff m = n \\ a <: \infty \end{array} \right\} \quad (\text{no finite age dereliction ; recall that } \uparrow^0 = \nu)$$

Rule `TY-TERM-PATU` is elimination (or pattern-matching) for unit, and is also used to chain destination-filling operations.

Pattern-matching with rules `TY-TERM-APP`, `TY-TERM-PATS`, `TY-TERM-PATP` and `TY-TERM-PATF` is parametrized by a mode m by which the typing context Θ_1 of the scrutinee is multiplied. The variables which bind the subcomponents of the scrutinee then inherit this mode. In particular, this choice crystalize the equivalence $!_{\omega a}(T_1 \otimes T_2) \simeq (!_{\omega a} T_1) \otimes (!_{\omega a} T_2)$, which is not part of intuitionistic linear logic, but valid in Linear Haskell [Bernardy et al. 2018]. We omit the mode annotation on **case** statements and lambda abstractions when the mode in question is the multiplicative neutral element 1ν of the mode semiring.

The Rule `TY-TERM-MAP` is where most of the safety of the system lies, and it is there where scope control takes place. It opens an ampar t , and binds its right side (containing destinations for holes on the other side, among other things) to variable x and then execute body t' . This lets the user access destinations of an ampar while temporarily forgetting about the structure with holes (being mutated behind the scenes by the destination-filling primitives). Type safety for **map** is based on the idea that a new scope is created for x and t' , so anything already present in the ambient scope (represented by Θ_2 in the conclusion) appears older when we see it from t' point of view. Indeed, when entering a new scope, the age of every remaining binding from the previous scopes is incremented by \uparrow . That way we can distinguish x from anything else that was already bound using the age of bindings alone. That's why t' types in $\uparrow\uparrow\Theta_2$, $x : \uparrow\nu$ while the global term **map** t with $x \mapsto t'$ types in Θ_1 , Θ_2 (notice the absence of shift on Θ_2). A schematic explanation of the scope rules is given in Figure 7.

We see in the figure that the left of an ampar (the structure being built) "takes place" in the ambient scope. The right side however, where destinations are, has its own new, inner scope that is opened when **mapped** over. When filling a destination (e.g. $x_1 \blacktriangleleft x_0$ in the figure), the right operand must be from a scope \uparrow older than the destination on the left of the operator, as this value will end up on the left of the ampar (which is thus in a scope \uparrow older than the destination originating from the right side).

(Yann) This sentence is incorrect, maybe it's missing a "what"? But I'm not sure, so I let you fix it

revisit if we allow weakening for dests, as we could replace that by $\omega\nu\Theta$ as in [Bernardy et al. 2018]

All the bit about mode ordering should move in the section about modes

"Crystalize"?

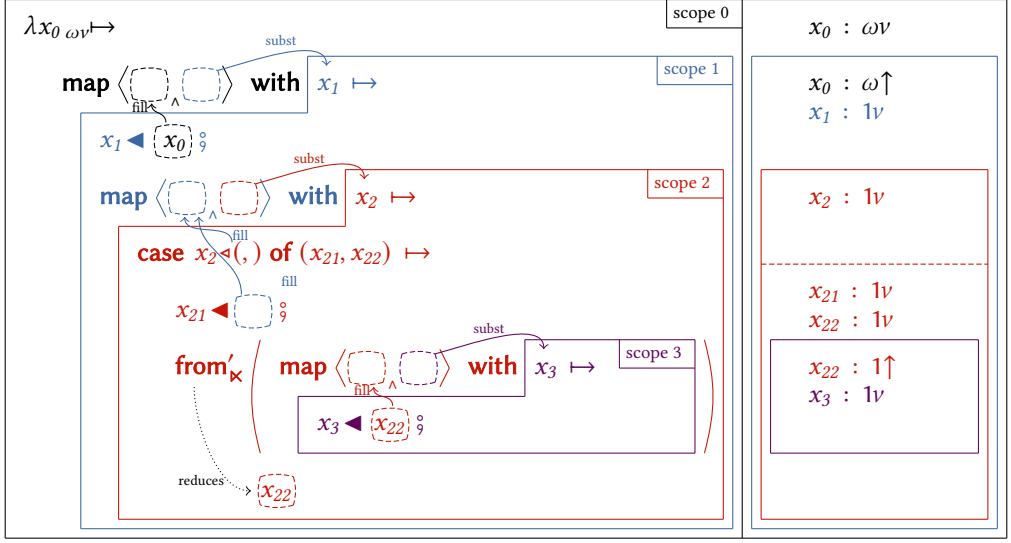
"Most of the safety" seems to imply that safety can be split into pieces. What you mean is that its the most interesting rule for destinations. Which also means that it should probably be explained much earlier.

You have never defined "open an ampar"

I'm not sure terms have a point of view

I didn't understand this paragraph

Modes in context

Fig. 7. Scope rules for **map** in λ_d

The rule **TY-TERM-FILLCOMP**, or its simpler variant, **TY-STERM-FILLLEAF** from Figure 6 confirm this intuition. The left operand of these operators must be a destination that types in the ambient context (both Θ_1 unchanged in the premise and conclusion of the rules). The right operand, however, is a value that types in a context Θ_2 in the premise, but requires $1\uparrow\Theta_2$ in the conclusion. This is the opposite of the shift that **map** does: while **map** opens a child scope for its body, “fillComp” (\Leftarrow)/“fillLeaf” (\blacktriangleleft) opens a portal to the parent scope for their right operand, as seen in the schema. The same phenomenon happens for the resources captured by the body of a lambda abstraction in **TY-TERM-FILLF**.

In **TY-TERM-TOA**, the operator **to_K** embeds an already completed structure in an ampar whose left side is the structure, and right side is unit.

When using **from_K** (rule **TY-STERM-FROMA'**), the left of an ampar is extracted to the ambient scope (as seen at the bottom of Figure 7 with x_{22}): this is the fundamental reason why the left of an ampar has to “take place” in the ambient scope. We know the structure is complete and can be extracted because the right side is of type unit (1), and thus no destination on the right side means no hole can remain on the left. **from_K** is implemented in terms of **from_K** in Figure 4 to keep the core calculus tidier (and limit the number of typing rules, evaluation contexts, etc), but it can be implemented much more efficiently in a real-world implementation.

When an ampar is complete and disposed of with the more general **from_K** in rule **TY-TERM-FROMA** however, we extract both sides of the ampar to the ambient scope, even though the right side is normally in a different scope. This is only safe to do because the right side is required to have type $!_{\infty}T$, which means it is scope-insensitive (it cannot contain any scope-controlled resource). This also ensures that the right side cannot contain destinations, meaning that the structure on the left is complete and ready to be read.

The remaining operators $\Leftarrow()$, $\Leftarrow!n$, $\Leftarrow!nr$, $\Leftarrow!m$, $\Leftarrow()$ from rules **TY-TERM-FILL*** are the destination-filling primitives. They write a hollow constructor to the hole pointed by the destination operand, and return the potential new destinations that are created in the process (or unit if there is none).

$\boxed{h} :_{lv} v : T$			(Typing judgment for values)		
TY-VAL-HOLE	TY-VAL-DEST	TY-VAL-UNIT	TY-VAL-FUN		
$\frac{}{\boxed{h} :_{lv} T \vdash \boxed{h} : T}$	$\frac{}{\rightarrow h :_{lv} [n] T \vdash \rightarrow h : [n] T}$	$\frac{}{\bullet \vdash () : 1}$	$\frac{\Delta, x :_{m} T \vdash u : U}{\Delta \vdash \lambda x_{m} \mapsto u : T_{m} \rightarrow U}$		
TY-VAL-LEFT	TY-VAL-RIGHT	TY-VAL-PROD	TY-VAL-EXP		
$\frac{\Gamma \vdash v_1 : T_1}{\Gamma \vdash \text{Inl } v_1 : T_1 \oplus T_2}$	$\frac{\Gamma \vdash v_2 : T_2}{\Gamma \vdash \text{Inr } v_2 : T_1 \oplus T_2}$	$\frac{\Gamma_1 \vdash v_1 : T_1 \quad \Gamma_2 \vdash v_2 : T_2}{\Gamma_1 + \Gamma_2 \vdash (v_1, v_2) : T_1 \otimes T_2}$	$\frac{\Gamma \vdash v' : T}{n\Gamma \vdash E_n v' : !n T}$		
<div> <div>TY-VAL-AMPAR</div> <div> <div>LinOnly Δ_3</div> <div>FinAgeOnly Δ_3</div> <div> $\uparrow \Delta_1, \Delta_3 \vdash v_1 : T$ </div> <div> $\Delta_2, (\rightarrow^{\uparrow} \Delta_3) \vdash v_2 : U$ </div> </div> <div> $\Delta_1, \Delta_2 \vdash \text{hnames}_{(\Delta_3)} \langle v_2 \wedge v_1 \rangle : U \ltimes T$ </div> </div>					

Fig. 8. Typing rules for values

5.3.2 Typing of values \vdash . **TODO: Re-introduce Ty-term-val**

The class of values v is purely a denotation for runtime data structures.

The typing of runtime values, given in Figure 8 is where hole and destination bindings appears. Values can have holes and destinations inside, but a value used as a term must not have any hole³. Also, variables cannot mention free variables; that makes it easier to prove substitution properties (as we will see in Section 6.3, we perform substitutions not only in terms, but also in evaluation contexts sometimes).

Hole bindings and destination bindings of the same hole name h are meant to annihilate each other in the typing context given they have a matching base type and mode. That way, the typing context of a term can stay constant during reduction:

- when destination-filling primitives are evaluated to build up data structures, they linearly consume a destination and write to a hole at the same time which makes both disappear, thus the typing context stays balanced;
- when a new hole is created, a matching destination is returned too, so the typing context stays balanced too.

However, the annihilation between a destination and a hole binding having the same name in the typing tree is only allowed to happen around an ampar, as it is the ampar connective that bind the name across the two sides (the names bound are actually stored in a set H on the ampar value $H \langle v_2 \wedge v_1 \rangle$). In fact, an ampar can be seen as a sort of lambda-abstraction, whose body (containing holes instead of variables) and input sink (destinations) are split across two sides, and magically interconnected through the ampar connective.

In most rules, we use a sum $\Gamma_1 + \Gamma_2$ for typing contexts (or the disjoint variant Γ_1, Γ_2). This sum doesn't not allow for annihilation of bindings with the same name; the operation is partial, and in particular it isn't defined if a same hole name is present in the operands in the two different forms (hole binding and destination binding). In particular, a pair $(\boxed{h}, \rightarrow h)$ is not well-typed. A single typing context Γ is not allowed either to contain both a hole binding and a destination binding for the same hole name.

Typing of ampars. As stated above, the core idea of TY-VAL-AMPAR is to act as a binding connective

³see TY-TERM-VAL: the value must type in a context Δ which means “destination only” by convention

“Annihilate” isn't an example of a definition that is under-derused. It's used only quite locally, but it's actually relied on. That being said, I still think it has to go because it's under-derused and will probably raise more question that it will answer.

Can we kill all the occurrences of the word “consume”? I'd like this word, to become standard, because it's a key element of linear types' semantics. But I don't think we have room for giving it a proper definition, so we should try working around

for holes and destinations.

We define a new operator \rightarrow^{\cdot} to represent the matching hole bindings for a set of destination bindings. It is a partial, point-wise operation on typing bindings of a context where:

$$\rightarrow^{\cdot}(\rightarrow h :_{\text{iv}} \lfloor_n T \rfloor) = \boxed{h} :_n T$$

Only an input context Δ made only of destination bindings, with mode $1v$, results in a valid output context (which is then only composed of hole bindings).

Equipped with this operator, we introduce the annihilation of bindings using Δ_3 to represent destinations on the right side v_1 , and $\rightarrow^{\cdot} \Delta_3$ to represent the matching hole bindings on the left side v_2 (the structure under construction). Those bindings are only present in the premises of the rule but get removed in the conclusion. Those hole names are only local, bound to that ampar and don't affect the outside world nor can be referenced from the outside.

Both sides of the ampar may also contain stored destinations from other scopes, represented by $\uparrow \Delta_1$ and Δ_2 in the respective typing contexts of v_1 and v_2 . All holes introduced by this ampar have to be annihilated by matching destinations; following our naming convention, no hole binding can appear in Δ_1 , Δ_2 in the conclusion.

The properties $\text{LinOnly } \Delta_3$ and $\text{FinAgeOnly } \Delta_3$ are true given that $\rightarrow^{\cdot} \Delta_3$ is a valid typing context, so are not really a new restriction on Δ_3 . They are mostly used to ease the mechanical proof of type safety for the system.

Other notable typing rules for values. Rules TY-VAL-HOLE and TY-VAL-DEST indicates that a hole or destination must have mode $1v$ in the typing context to be used (except when a destination is stored away, as we have seen).

Rules for unit, left and right variants, and product are straightforward.

Rule TY-VAL-EXP is rather classic too: we multiply the dependencies Γ of the value by the mode n of the exponential. The intuition is that if v uses a resource v' twice, then $e_2 v$, that corresponds to two uses of v (in a system with such a mode), will use v' four times.

Rule TY-VAL-FUN indicates that (value level) lambda abstractions cannot have holes inside. In other terms, a function value cannot be built piecemeal like other data structures, its whole body must be a complete term right from the beginning. It cannot contain free variables either, as the body of the function must type in context Δ , $x :_m T$ where Δ is made only of destination bindings. One might wonder, how can we represent a curried function $\lambda x \mapsto \lambda y \mapsto x \text{ concat } y$ as the value level, as the inner abstraction captures the free variable x ? The answer is that such a function, at value level, is encoded as $\lambda x \mapsto \text{from}'_x (\text{map alloc with } d \mapsto d \triangleleft (\lambda y \mapsto x \text{ concat } y))$, where the inner closure is not yet in value form, but pending to be built into a value. As the form $d \triangleleft (\lambda y \mapsto t)$ is part of term syntax, and not value syntax, we allow free variable captures in it.

6 EVALUATION CONTEXTS AND SEMANTICS

The semantics of λ_d is given using small-step reductions on a pair $C[t]$ of an evaluation context C (represented by a stack) determining a focusing path, and a term t under focus. Such a pair $C[t]$ is called a *command*, and represents a running program. We think that our semantics makes it easier to reason about types and linearity of a running program with holes than a store-based approach such as the one previewed in Section 3.

6.1 Evaluation contexts forms

The grammar of evaluation contexts is given in Figure 9. An evaluation context C is the composition of an arbitrary number of focusing components $c_1, c_2 \dots$. We chose to represent this composition explicitly using a stack, instead of a meta-operation that would only let us access its final result. As

```

785  c ::= t' [] | [] v | [] § u
786      | casem [] of {lnl x1 ↦ u1, lnr x2 ↦ u2} | casem [] of (x1, x2) ↦ u | casem [] of En x ↦ u
787      | map [] of x ↦ t' | toκ [] | fromκ []
788      | [] < () | [] < lnl | [] < lnr | [] < (.) | [] < Em | [] < (λxm ↦ u) | [] < o t' | v < o []
789      | opH(v2 ∧ []) (open ampar focus component)
790
791  C ::= [] | C o c | C[h:=H v]

```

Fig. 9. Grammar for evaluation contexts

a result, focusing and defocusing operations are made explicit in the semantics, resulting in a more verbose but simpler proof. It is also easier to imagine how to build a stack-based interpreter for such a language.

Focusing components are all directly derived from the term syntax, except for the “open ampar” focus component $\overset{\text{op}}{H}(v_2 \wedge [])$. This focus component indicates that an ampar is currently being mapped on, with its left-hand side v_2 (the structure being built) being attached to the “open ampar” focus component, while its right-hand side (containing destinations) is either in subsequent focus components, or in the term under focus.

We introduce a special substitution $C[h:=_H v]$ that is used to update structures under construction that are attached to open ampar focus components in the stack. Such a substitution is triggered when a destination $\rightarrow h$ is filled in the term under focus, and results in the value v (that may contain holes itself, e.g. if it is a hollow constructor $(\boxed{h_1}, \boxed{h_2})$) being written to the hole \boxed{h} (that must appear somewhere on an open ampar payload). The set H tracks the potential hole names introduced by value v .

6.2 Typing of evaluation contexts and commands

Evaluation contexts are typed in a context Δ that can only contains destination bindings. Δ represents the typing context available for the term that will be put in the box $[]$ of the evaluation context. In other terms, while the typing context of a term is a list of requirements so that it can be typed, the typing context of an evaluation context is the set of bindings that it makes available to the term under focus. As a result, while the typing of a term $\Theta \vdash t : T$ is additive (the typing requirements for a function application is the sum of the requirements for the function itself and for its argument), the typing of an evaluation context $\Delta \dashv C : T \rightarrow U_0$ is subtractive : adding the focus component $t' []$ to the stack C will remove whatever is needed to type t' from the typing context provided by C . The whole typing rules for evaluation contexts C as well as commands $C[t]$ are presented in Figure 10.

An evaluation context has a pseudo-type $T \rightarrow U_0$, where T denotes the type of the focus (i.e. the type of the term that can be put in the box of the evaluation context) while U_0 denotes the type of the resulting command (when the box of the evaluation context is filled with a term).

Composing an evaluation context of pseudo-type $T \rightarrow U_0$ with a new focus component never affects the type U_0 of the future command ; only the type T of what can be put in the box is altered.

All typing rules for evaluation contexts can be derived from the ones for the corresponding term (except for the rule TY-ECTXS-OPENAMPAR that is the truly new form). Let's take the rule TY-ECTXS-PATP as an example:

- the typing context $m \Delta_1 + \Delta_2$ in the premise for C corresponds to $m \Theta_1 + \Theta_2$ in the conclusion of TY-TERM-PATP in Figure 6;

“Box” is a very atypical term for this. Usually it's a “hole” but we've been using it for something else. Can we find a better way to say this sentence? Below we use “focus” for it, it's not a completely orthodox use of the term either, but it's better.

It's hard to understand what's going on here. Let's spell out the form of the judgment in the sentence, and point at bits. Also, prefer saying what things are before saying what they're not rather than

834	$\boxed{\Delta \vdash C : T \rightarrow U_0}$		(Typing judgment for evaluation contexts)
835		TY-ECTXS-APP1	TY-ECTXS-APP2
836		$m\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$	$m\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$
837	TY-ECTXS-ID	$\Delta_2 \vdash t' : T_m \rightarrow U$	$\Delta_1 \vdash v : T$
838	$\vdash [] : U_0 \rightarrow U_0$	$\Delta_1 \vdash C \circ (t' []) : T \rightarrow U_0$	$\Delta_2 \vdash C \circ ([\] \vee) : (T_m \rightarrow U) \rightarrow U_0$
839		TY-ECTXS-PAT	
840			
841	TY-ECTXS-PATU		$m\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$
842	$\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$		$\Delta_2, x_1 : mT_1 \vdash u_1 : U$
843	$\Delta_2 \vdash u : U$		$\Delta_2, x_2 : mT_2 \vdash u_2 : U$
844	$\Delta_1 \vdash C \circ ([\] \circ u) : 1 \rightarrow U_0$	$\Delta_1 \vdash C \circ (\text{case}_m [\] \text{ of } \{ \text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2 \}) : (T_1 \oplus T_2) \rightarrow U_0$	
845	TY-ECTXS-PATP		TY-ECTXS-PATE
846	$m\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$		$m\Delta_1, \Delta_2 \vdash C : U \rightarrow U_0$
847	$\Delta_2, x_1 : mT_1, x_2 : mT_2 \vdash u : U$		$\Delta_2, x : m-m' T \vdash u : U$
848	$\Delta_1 \vdash C \circ (\text{case}_m [\] \text{ of } (x_1, x_2) \mapsto u) : (T_1 \otimes T_2) \rightarrow U_0$		$\Delta_1 \vdash C \circ (\text{case}_m [\] \text{ of } E_{m'} x \mapsto u) : !_{m'} T \rightarrow U_0$
849			
850	TY-ECTXS-MAP		TY-ECTXS-TOA
851	$\Delta_1, \Delta_2 \vdash C : U \times T' \rightarrow U_0$		$\Delta \vdash C : (U \times 1) \rightarrow U_0$
852	$\uparrow \Delta_2, x : !_{\text{iv}} T \vdash t' : T'$		$\Delta \vdash C \circ (\text{to}_\times [\]) : U \rightarrow U_0$
853	$\Delta_1 \vdash C \circ (\text{map} [\] \text{ of } x \mapsto t') : (U \times T) \rightarrow U_0$		
854	TY-ECTXS-FROMA		TY-ECTXS-FILLU
855	$\Delta \vdash C : (U \otimes (!_{100} T)) \rightarrow U_0$		$\Delta \vdash C : 1 \rightarrow U_0$
856	$\Delta \vdash C \circ (\text{from}_\times [\]) : (U \times (!_{100} T)) \rightarrow U_0$		$\Delta \vdash C \circ ([\] \triangleleft ()) : [n] 1 \rightarrow U_0$
857			
858	TY-ECTXS-FILLL		TY-ECTXS-FILLR
859	$\Delta \vdash C : [n] T_1 \rightarrow U_0$		$\Delta \vdash C : [n] T_2 \rightarrow U_0$
860	$\Delta \vdash C \circ ([\] \triangleleft \text{Inl}) : [n] T_1 \oplus T_2 \rightarrow U_0$		$\Delta \vdash C \circ ([\] \triangleleft \text{Inr}) : [n] T_1 \oplus T_2 \rightarrow U_0$
861			
862	TY-ECTXS-FILLP		TY-ECTXS-FILLE
863	$\Delta \vdash C : ([n] T_1 \otimes [n] T_2) \rightarrow U_0$		$\Delta \vdash C : [m-n] T \rightarrow U_0$
864	$\Delta \vdash C \circ ([\] \triangleleft (,)) : [n] T_1 \otimes T_2 \rightarrow U_0$		$\Delta \vdash C \circ ([\] \triangleleft E_m) : [n] !_{m'} T \rightarrow U_0$
865			
866	TY-ECTXS-FILLF		TY-ECTXS-FILLCOMP1
867	$\Delta_1, (\uparrow \cdot n) \Delta_2 \vdash C : 1 \rightarrow U_0$		$\Delta_1, (\uparrow \cdot n) \Delta_2 \vdash C : T \rightarrow U_0$
868	$\Delta_2, x : mT \vdash u : U$		$\Delta_2 \vdash t' : U \times T$
869	$\Delta_1 \vdash C \circ ([\] \triangleleft (\lambda x_m \mapsto u)) : [n] T_m \rightarrow U \rightarrow U_0$		$\Delta_1 \vdash C \circ ([\] \triangleleft t') : [n] U \rightarrow U_0$
870			
871		TY-ECTXS-OPENAMPAR	
872	TY-ECTXS-FILLCOMP2	$\text{LinOnly } \Delta_3 \quad \text{FinAgeOnly } \Delta_3$	
873	$\Delta_1, (\uparrow \cdot n) \Delta_2 \vdash C : T \rightarrow U_0$	$\text{hnames}(C) \quad \# \# \quad \text{hnames}(\Delta_3)$	
874	$\Delta_1 \vdash v : [n] U$	$\Delta_1, \Delta_2 \vdash C : (U \times T') \rightarrow U_0$	
875	$\Delta_2 \vdash C \circ (v \triangleleft [\]) : U \times T \rightarrow U_0$	$\Delta_2, \rightarrow \cdot \Delta_3 \quad \forall \vdash v_2 : U$	
876	$\boxed{\vdash C[t] : T}$	$\uparrow \Delta_1, \Delta_3 \vdash C \circ (\overset{\text{op}}{\text{hnames}}(\Delta_3) \langle v_2 \wedge [\] \rangle) : T' \rightarrow U_0$	
877			(Typing judgment for commands)
878		TY-CMD	
879		$\Delta \vdash C : T \rightarrow U_0 \quad \Delta \vdash t : T$	
880		$\vdash C[t] : U_0$	
881			
882			

Fig. 10. Typing rules for evaluation contexts and commands

- the typing context $\Delta_2, x_1 :_{\mathbf{m}} T_1, x_2 :_{\mathbf{m}} T_2$ in the premise for term u corresponds to the typing context $\Theta_2, x_1 :_{\mathbf{m}} T_1, x_2 :_{\mathbf{m}} T_2$ for the same term in TY-TERM-PATP;
- the typing context Δ_1 in the conclusion for $C \circ (\text{case}_{\mathbf{m}} [] \text{ of } (x_1, x_2) \mapsto u)$ corresponds to the typing context Θ_1 in the premise for t in TY-TERM-PATP (the term t is located where the box $[]$ is in TY-ECTXS-OPENAMPAR).

In a way, the typing rule for an evaluation context is a “rotation” of the typing rule for the associated term, where the typing contexts of one premise and the conclusion are swapped, and the typing context of the other potential premise is kept unchanged (with the added difference that free variables cannot appear in typing contexts of evaluation contexts, so any Θ becomes a Δ).

As we see at the bottom of the figure, a command $C[t]$ (i.e. a pair of an evaluation context and a term) is well typed when the evaluation context C provides a typing context Δ that is exactly one in which t is well typed. We can always embed a well-typed, closed term $\bullet \vdash t : T$ as a well-typed command using the identity evaluation context: $t \simeq [][t]$ and we thus have $\vdash [][t] : T$ where $\Delta = \bullet$ (the empty context).

6.3 Small-step semantics

We equip λ_d with small-step semantics. There are three types of semantic rules:

- focus rules, where we remove a layer from non-value term t and push a corresponding focus component on the stack C ;
- unfocus rules, where t is a value and thus we pop a focus component from the stack C and transform it back to a term, so that a redex appears (or so that another focus/unfocus rule can be triggered);
- reduction rules, where the actual computation logic takes place.

Here is the whole set of rules for PATP, in the previously announced order:

$$\begin{aligned}
 C[\text{case}_{\mathbf{m}} t \text{ of } (x_1, x_2) \mapsto u] &\longrightarrow (C \circ (\text{case}_{\mathbf{m}} [] \text{ of } (x_1, x_2) \mapsto u))[t] \quad \text{when } \text{NotVal } t \\
 (C \circ (\text{case}_{\mathbf{m}} [] \text{ of } (x_1, x_2) \mapsto u))[v] &\longrightarrow C[\text{case}_{\mathbf{m}} v \text{ of } (x_1, x_2) \mapsto u] \\
 C[\text{case}_{\mathbf{m}} (v_1, v_2) \text{ of } (x_1, x_2) \mapsto u] &\longrightarrow C[u[x_1 := v_1][x_2 := v_2]]
 \end{aligned}$$

Rules are triggered in a purely deterministic fashion; once a subterm is a value, it cannot be focused again. As focusing and defocusing rules are entirely mechanical (they are just a matter of pushing and popping a focus component on the stack), we only present the set of reduction rules for the system in Figure 11, but the whole system is included in the annex (see Figures 12 and 13).

Reduction rules for function application, pattern-matching, **to**_κ and **from**_κ are straightforward.

All reduction rules for destination-filling primitives trigger a substitution $C[h :=_H v]$ on the evaluation context C that corresponds to a memory update of a hole $[h]$. RED-FILLU and RED-FILLF do not create any new hole; they only write a value to an existing one. On the other hand, rules RED-FILLL, RED-FILLR, RED-FILLE and RED-FILLP all write a hollow constructor to the hole h , that is to say a value containing holes itself. Thus, we need to generate fresh names for these new holes, and also return a destination for each new hole with a matching name.

Obtaining a fresh name is represented by the statement $h' = \max(\text{hnames}(C) \cup \{h\}) + 1$ in the premises of these rules. One invariant of the system is that an ampar must have fresh names to be opened, so we always rename local hole names bound by an ampar to fresh names just when that ampar is mapped on, as these local names — represented by the set of hole names H that the ampar carries — could otherwise shadow already existing names in evaluation context C . This invariant is materialized by premise $\text{hnames}(C) \# \# \text{hnames}(\Delta_3)$ in rule TY-ECTXS-OPENAMPAR for the open ampar focus component that is created during reduction of a **map**.

We use hole name shifting as a strategy to obtain fresh names. Shifting all hole names in a set H

“in a way” AND “rotation” between quotes. This is a little too much analogy for a text like this. Let’s keep at something like “The typing rule for an evaluation context can be derived systematically from the typing rules of the corresponding term form”

This paragraph is difficult to understand

I’m not sure “layer” is the appropriate term here

For some reason I’m only noticing the name “focus component” now. There’s nothing really wrong with it, but I’m not fully convinced either.

The “previously announced order”? Did we announce an order previously?

Arnaud: ah yes, you (Thomas) told me of this I think. The fact that the notation for the focus and for hole-substitution are so similar is a bit confusing indeed.

For substitution to “correspond” to

932	$\boxed{C[t]} \longrightarrow C'[t']$	(Small-step evaluation of commands)
933	$C[(\lambda x_m \mapsto u) v] \longrightarrow C[u[x := v]]$	RED-APP
934	$C[() \circ u] \longrightarrow C[u]$	RED-PATU
935	$C[\text{case}_m (\text{Inl } v_1) \text{ of } \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}] \longrightarrow C[u_1[x_1 := v_1]]$	RED-PATL
936	$C[\text{case}_m (\text{Inr } v_2) \text{ of } \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}] \longrightarrow C[u_2[x_2 := v_2]]$	RED-PATR
937	$C[\text{case}_m (v_1, v_2) \text{ of } (x_1, x_2) \mapsto u] \longrightarrow C[u[x_1 := v_1][x_2 := v_2]]$	RED-PATP
938	$C[\text{case}_m E_n v' \text{ of } E_n x \mapsto u] \longrightarrow C[u[x := v']]$	RED-PATE
939	$C[\text{to}_K v_2] \longrightarrow C[\{\} \langle v_2 \wedge () \rangle]$	RED-TOA
940	$C[\text{from}_K \{\} \langle v_2 \wedge E_{100} v_1 \rangle] \longrightarrow C[(v_2, E_{100} v_1)]$	RED-FROMA
941	$C[\rightarrow h \triangleleft ()] \longrightarrow C[h := \{\} ()][()]$	RED-FILLU
942	$C[\rightarrow h \triangleleft (\lambda x_m \mapsto u)] \longrightarrow C[h := \{\} \lambda x_m \mapsto u][()]$	RED-FILLF
943	$C[\rightarrow h \triangleleft \text{Inl}] \longrightarrow C[h := \{h'+1\} \text{Inl } \boxed{h'+1}][\rightarrow h'+1]$	RED-FILLL
944	$C[\rightarrow h \triangleleft \text{Inr}] \longrightarrow C[h := \{h'+1\} \text{Inr } \boxed{h'+1}][\rightarrow h'+1]$	RED-FILLR
945	$C[\rightarrow h \triangleleft E_m] \longrightarrow C[h := \{h'+1\} E_m \boxed{h'+1}][\rightarrow h'+1]$	RED-FILLE
946	$C[\rightarrow h \triangleleft ()] \longrightarrow C[h := \{h'+1, h'+2\} (\boxed{h'+1}, \boxed{h'+2})][(\rightarrow h'+1, \rightarrow h'+2)]$	RED-FILLP
947	$C[\rightarrow h \triangleleft H \langle v_2 \wedge v_1 \rangle] \longrightarrow C[h := (H \pm h') v_2 [H \pm h']][v_1 [H \pm h']]$	RED-FILLCOMP
948	$C[\text{map } H \langle v_2 \wedge v_1 \rangle \text{ with } x \mapsto t'] \longrightarrow (C \circ (H \pm h')^{\text{op}} \langle v_2 [H \pm h''] \wedge [] \rangle)[t' [x := v_1 [H \pm h'']]]$	OPEN-AMPAR
949	$(C \circ H^{\text{op}} \langle v_2 \wedge [] \rangle)[v_1] \longrightarrow C[H \langle v_2 \wedge v_1 \rangle]$	CLOSE-AMPAR
950	$\text{where } \begin{cases} h' &= \max(\text{hnames}(C) \cup \{h\}) + 1 \\ h'' &= \max(\text{hnames}(C)) + 1 \end{cases}$	

Fig. 11. Small-step semantics

by a given offset h' is denoted $H \pm h'$. We extend this notation to define a conditional shift operation $\boxed{H \pm h'}$ which shifts each hole name appearing in the operand to the left of the brackets by h' if this hole name is also member of H . This conditional shift can be used on a single hole name, a value, or a typing context.

In rule RED-FILLCOMP, we write the left-hand side v_2 of a closed ampar $H \langle v_2 \wedge v_1 \rangle$ to a hole \boxed{h} that is part of some focus fragment $H' \langle v_2' \wedge [] \rangle$ in the evaluation context C . That fragment is not mentioned explicitly in the rule, as the destination $\rightarrow h$ is enough to target it. This results in the composition of two structures with holes v_2' and v_2 through filling of $\rightarrow h$. Because we split open the ampar $H \langle v_2 \wedge v_1 \rangle$ (its left-hand side gets written to a hole, while its right hand side is returned), we need to rename any hole name that it contains to a fresh one, as we do when an ampar is opened in the **map** rule. The renaming is carried out by the conditional shift $v_2 [H \pm h']$ and $v_1 [H \pm h']$ (only hole names local to the ampar, represented by the set H , gets renamed).

Last but not least, rules OPEN-AMPAR and CLOSE-AMPAR dictates how and when a closed ampar (a term) is converted to an open ampar (a focusing fragment) and vice-versa. With OPEN-AMPAR, the local hole names of the ampar gets renamed to fresh ones, and the left-hand side gets attached to the focusing fragment $H \pm h' \langle v_2 [H \pm h'] \wedge [] \rangle$ while the right-hand side (containing destinations) is substituted in the body of the **map** statement (which becomes the new term under focus). This effectively allows the right-hand side of an ampar to be a term instead of a value for a limited time.

The rule CLOSE-AMPAR triggers when the body of a **map** statement has reduced to a value. In that case, we can close the ampar, by popping the focus fragment from the stack C and merging back with v_2 to reform a closed ampar.

Type safety. With the semantics now defined, we can state the usual type safety theorems:

THEOREM 6.1 (TYPE PRESERVATION). *If $\vdash C[t] : T$ and $C[t] \longrightarrow C'[t']$ then $\vdash C'[t'] : T$.*

THEOREM 6.2 (PROGRESS). *If $\vdash C[t] : T$ and $\forall v, C[t] \neq [][v]$ then $\exists C', t'. C[t] \longrightarrow C'[t']$.*

A command of the form $[][v]$ cannot be reduced further, as it only contains a fully determined value, and no pending computation. This is the stopping point of the reduction, and any well-typed command eventually reaches this form.

7 FORMAL PROOF OF TYPE SAFETY

We've proved type preservation and progress theorems with the Coq proof assistant. At time of writing, we have assumed, rather than proved, the substitution lemmas. The choice of turning to a proof assistant was a pragmatic choice: the context handling in λ_d can be quite finicky, and it was hard, without computer assistance, to make sure that we hadn't made mistakes in our proofs. The version of λ_d that we've proved is written in Ott [Sewell et al. 2007], the same Ott file is used as a source for this article, making sure that we've proved the same system as we're presenting; some visual simplification is applied by a script to produce the version in the article.

Most of the proof was done by an author with little prior experience with Coq. This goes to show that Coq is reasonably approachable even for non-trivial development. The proof is about 6000 lines long, and contains nearly 350 lemmas. Many of the cases of the type preservation and progress lemmas are similar. To handle such repetitive cases, the use of a large-language-model based autocompletion system has proven quite effective.

Binders are the biggest problem. We've largely manage to make the proof to be only about closed terms, to avoid any complication with binders. This worked up until the substitution lemmas, which is the reason why we haven't proved them in Coq yet (that and the fact that it's much easier to be confident in our pen-and-paper proofs for those). There are backends to generate locally nameless representations from Ott definitions [Nardelli 2009]; we haven't tried them yet, but the unusual binding nature of ampars may be too much for them to handle.

The proofs aren't very elegant. For instance, we don't have any abstract formalization of semirings: since our semirings are finite it was more expedient to brute-force the properties we needed by hand. We've observed up to 232 simultaneous goals, but a computer makes short work of this: it was solved by a single call to the congruence tactic. Nevertheless there are a few points of interest:

- we represent context as finite-domain functions, rather than as syntactic lists. This works much better when defining sums of context. There are a bunch of finite-function libraries in the ecosystem, but we needed finite dependent functions (because the type of binders depend on whether we're binding a variable name or a hole name). This didn't exist, but for our limited purpose, it ended up not being too costly rolling our own. About 1000 lines of proofs. The underlying data type is actual functions, this was simpler to develop, but equality is more complex than with a bespoke data type;
- addition of context is partial since we can only add two binding of the same name if they also have the same type. Instead of representing addition as a binary function to an optional context, we represent addition as a total function to contexts, but we change contexts to have faulty bindings on some names. This simplifies reasoning about properties like commutativity and associativity, at the cost of having well-formedness preconditions in the premises of typing rules as well as some lemmas.

The inference rules produced by Ott aren't conducive to using setoid equality. This turned out to be a problem with our type for finite function:

```
Record T A B := {
  underlying :> forall x:A, option (B x);
  supported : exists l : list A, Support l underlying;
}.
```

where `Support l f` means that `l` contains the domain of `f`. To make the equality of finite function be strict equality `eq`, we assumed functional extensionality and proof irrelevance. In some circumstances, we've also needed to list the finite functions' domains. But in the definition, the domain is sealed behind a proposition, so we also assumed classical logic as well as indefinite description

```
Axiom constructive_indefinite_description :
  forall (A : Type) (P : A->Prop), (exists x, P x) -> { x : A | P x }.
```

together, they let us extract the domain from the proposition. Again this isn't particularly elegant, we could have avoided some of these axioms at the price of more complex development. But for the sake of this article, we decided to favor expediency over elegance.

8 IMPLEMENTATION OF λ_d USING IN-PLACE MEMORY MUTATIONS

The formal language presented in Sections 5 and 6 is not meant to be implemented as-is.

First, λ_d misses a form of recursion, but we believe that adding equirecursive types and a fix-point operator wouldn't compromise the safety of the system.

Secondly, ampars are not managed linearly in λ_d ; only destinations are. That is to say that an ampar can be wrapped in an exponential, e.g. $E_{\omega V} \{h\} \langle 0 :: \boxed{h} \rightarrow h \rangle$ (representing a non-linear difference list $0 :: \boxed{}$), and then used twice, each time in a different way:

```
case  $E_{\omega V} \{h\} \langle 0 :: \boxed{h} \rightarrow h \rangle$  of  $E_{\omega V} x \mapsto$ 
  let  $x_1 := x$  append 1 in
  let  $x_2 := x$  append 2 in
  toList ( $x_1$  concat  $x_2$ )
 $\longrightarrow^* 0 :: 1 :: 0 :: 2 :: []$ 
```

It may seem counter-intuitive at first, but this program is valid and safe in λ_d . Thanks to the renaming discipline we detailed in Section 6.3, every time an ampar is **mapped** over, its hole names are renamed to fresh ones. So when we call **append** to build x_1 (which is implemented in terms of **map**), we sort of allocate a new copy of the ampar before mutating it, effectively achieving a copy-on-write memory scheme. Thus it is safe to operate on x again to build x_2 .

In the introduction of the article, we announced a safe framework for in-place memory mutation, so we will uphold this promise now. The key to go from a copy-on-write scheme to an in-place mutation scheme is to force ampars to be linearly managed too. For that we introduce a new type **Token**, together with primitives **dup** and **drop** (remember that unqualified arrows have mode **1v**, so are linear):

```
dup : Token  $\rightarrow$  Token@Token
drop : Token  $\rightarrow$  1
alloccow : T  $\times$  [T]
allocip : Token  $\rightarrow$  T  $\times$  [T]
```

We now have two possible versions of **alloc**: the new one with an in-place mutation memory model (**ip**), that has to be managed linearly, and the old one that doesn't have to be used linearly, and features a copy-on-write (cow) memory model.

We use the `Token` type as an intrinsic source of linearity that infects the ampar returned by `allocp`. Such a token can be duplicated using `dup`, but as soon as it is used to create an ampar, that ampar cannot be duplicated itself. In the system featuring the `Token` type and `allocp`, “closed” programs now typecheck in the non-empty context $\{tok_0 :_{100} \text{Token}\}$ containing a token variable that the user can `duplicate` and `drop` freely to give birth to an arbitrary number of ampars, that will then have to be managed linearly.

Having closed programs to typecheck in non-empty context $\{tok_0 :_{100} \text{Token}\}$ is very similar to having a primitive function `withToken` : $(\text{Token } 1_{100} \rightarrow !_{\omega_{100}} T) \rightarrow !_{\omega_{100}} T$ as it is done in [Bagrel 2024].

In such an extension, as ampars are managed linearly, we can change the allocation and renaming mechanisms:

- the hole name for a new ampar can be chosen fresh right from the start (this corresponds to a new heap allocation);
- adding a new hollow constructor still require freshness for its hole names (this corresponds to a new heap allocation too);
- `mapping` over an ampar and filling destinations or composing two ampars using “`fillComp`” (\Leftarrow) no longer require any renaming: we have the guarantee that the names are globally fresh, and thus we can do in-place memory updates.

We decided to omit the linearity aspect of ampars in λ_d as it clearly obfuscate the presentation of the system without adding much to the understanding of the latter. We believe that the system is still sound with this linearity aspect, and articles such as [Spiwack et al. 2022] gives a pretty clear view on how to implement the linearity requirement for ampars in practice without too much noise for the user.

9 RELATED WORK

9.1 Destination-passing style for efficient memory management

In [Shaikhha et al. 2017], the authors present a destination-based intermediate language for a functional array programming language. They develop a system of destination-specific optimizations and boast near-C performance.

This is the most comprehensive evidence to date of the benefit of destination-passing style for performance in functional programming languages. Although their work is on array programming, while this article focuses on linked data structure. They can therefore benefit of optimizations that are perhaps less valuable for us, such as allocating one contiguous memory chunk for several arrays.

The main difference between their work and ours is that their language is solely an intermediate language: it would be unsound to program in it manually. We, on the other hand, are proposing a type system to make it sound for the programmer to program directly with destinations.

We consider that these two aspects complement each other: good compiler optimization are important to alleviate the burden from the programmer and allowing high-level abstraction; having the possibility to use destinations in code affords the programmer more control would they need it.

9.2 Tail modulo constructor

Another example of destinations in a compiler’s optimizer is [Bour et al. 2021]. It’s meant to address the perennial problem that the map function on linked lists isn’t tail-recursive, hence consumes stack space. The observation is that there’s a systematic transformation of functions where the only recursive call is under a constructor to a destination-passing tail-recursive implementation.

(Yann) This part wasn’t entirely clear to me initially. I think what get me to understand are the following simple facts, that you might or might not want to stress more explicitly: 1. we can’t create tokens in the language, anywhere 2. Which means that you always need to have the token in the context, and in particular in the context of any subterm that contains an call to ‘alloc_p’ 3. Which means that any application of alloc_p creating an ampar will be contaminated with the (linear) token in the context and can’t be duplicated

Subsection for each work seems a little heavy-weight. On the other hand the italic paragraphs seem too light-weight. I’m longing for the days of paragraph heading in bold.

(Yann) This could be nuanced: allocating deeply nested trees like ASTs with a bump allocator so that they

Here again, there's no destination in user land, only in the intermediate representation. However, there is a programmatic interface: the programmer annotates a function like

```
let[@tail_mod_cons] rec map =
```

to ask the compiler to perform the translation. The compiler will then throw an error if it can't. This way, contrary to the optimizations in [Shaikhha et al. 2017], this optimization is entirely predictable.

This has been available in OCaml since version 4.14. This is the one example we know of of destinations built in a production-grade compiler. Our λ_d makes it possible to express the result tail-modulo-constructor in a typed language. It can be used to write programs directly in that style, or it could serve as a typed target language for and automatic transformation. On the flip-side, tail modulo constructor is too weak to handle our difference lists or breadth-first traversal examples.

TODO: Mention Tail modulo context

9.3 A functional representation of data structures with a hole

The idea of using linear types to let the user manipulate structures with holes safely dates back to [Minamide 1998]. Our system is strongly inspired by theirs. In their system, we can only compose functions that represent data structures with holes, we can't pattern-match on the result; just like in our system we cannot act on the left-hand side of $S \times T$, only the right hand part.

In [Minamide 1998], it's only ever possible to represent structures with a single hole. But this is a rather superficial restriction. The author doesn't comment on this, but we believe that this restriction only exists for convenience of the exposition: the language is lowered to a language without function abstraction and where composition is performed by combinators. While it's easy to write a combinator for single-argument-function composition, it's cumbersome to write combinators for functions with multiple arguments. But having multiple-hole data structures wouldn't have changed their system in any profound way.

The more important difference is that while their system is based on a type of linear functions, our is based on the linear logic's par combinator. This, in turns, lets us define a type of destinations which are representations of holes in values, which [Minamide 1998] doesn't have. This means that [Minamide 1998] can implement our examples with difference lists and queues from Section 2.2, but it can't do our breadth-first traversal example from Section 4, since storing destinations in a data structure is the essential ingredient of this example.

This ability to store destination does come at a cost though: the system needs this additional notion of ages to ensure that destinations are use soundly. On the other hand, our system is strictly more general, in that the system from [Minamide 1998] can be embedded in λ_d , and if one stays in this fragment, we're never confronted with ages. Ages only show up when writing programs that go beyond Minamide's system.

9.4 Destination-passing style programming: a Haskell implementation

In [Bagrel 2024], the author proposes a system much like ours: it has a destination type, and a par-like construct (that they call *Incomplete*), where only the right-hand side can be modified; together these two elements give extra expressiveness to the language compared to [Minamide 1998].

In their system, $d \blacktriangleleft t$ requires t to be unrestricted, while in λ_d , t can be linear. The consequence is that in [Bagrel 2024], destinations can be stored in data structures but not in data structures with holes. In order to do a breadth-first search algorithm like in Section 4, they can't use improved queues like we do, they have to use regular functional queues.

However, [Bagrel 2024] is implemented in Haskell, which just features linear types. Our system subsumes theirs; but it requires the age system that is more than what Haskell provides. Encoding their system in ours will unfortunately make ages appear in the typing rules.

9.5 Semi-axiomatic sequent calculus

In [DeYoung et al. 2020], the author develop a system where constructors return to a destination rather than allocating memory. It is very unlike the other systems described in this section in that it's completely founded in the Curry-Howard isomorphism. Specifically it gives an interpretation of a sequent calculus which mixes Gentzen-style deduction rules and Hilbert-style axioms. As a consequence, the par connective is completely symmetric, and, unlike our $[T]$ type, their dualization connective is involutive.

The cost of this elegance is that computations may try to pattern-match on a hole, in which case they must wait for the hole to be filled. So the semantic of holes is that of a future or a promise. In turns this requires the semantic of their calculus to be fully concurrent. Which is a very different point in the design space.

10 CONCLUSION AND FUTURE WORK

Using a system of ages in addition to linearity, λ_d is a purely functional calculus which supports destinations in a very flexible way. It subsumes existing calculi from the literature for destination passing, allowing both composition of data structures with holes and storing destinations in data structures. Data structures are allowed to have multiple holes, and destinations can be stored in data structures that, themselves, have holes. The latter is the main reason to introduce ages and is key to λ_d 's flexibility.

We don't anticipate that a system of ages like λ_d will actually be used in a programming language: it's unlikely that destination are so central to the design of a programming language that it's worth baking them so deeply in the type system. Perhaps a compiler that makes heavy use of destinations in its optimizer could use λ_d as a typed intermediate representation. But, more realistically, our expectation is that λ_d can be used as a theoretical framework to analyze destination-passing systems: if an API can be defined in λ_d then it's sound.

In fact, we plan to use this very strategy to design an API for destination passing in Haskell, leveraging only the existing linear types, but retaining the possibility of storing destinations in data structures with holes.

Missing here: in λ_d , destination-filling is complete as illustrate in Figure 4, whereas in Bagrel 2024, we need constructors (since those are linear and destination-filling isn't).

REFERENCES

- Andreas Abel and Jean-Philippe Bernardy. 2020. A unified view of modalities in type systems. *Proc. ACM Program. Lang.* 4, ICFP, Article 90 (aug 2020), 28 pages. <https://doi.org/10.1145/3408972>
- Thomas Bagrel. 2024. Destination-passing style programming: a Haskell implementation. In *35es Journées Francophones des Langages Applicatifs (JFLA 2024)*. Saint-Jacut-de-la-Mer, France. <https://inria.hal.science/hal-04406360>
- Jean-Philippe Bernardy, Mathieu Boespflug, Ryan R. Newton, Simon Peyton Jones, and Arnaud Spiwack. 2018. Linear Haskell: practical linearity in a higher-order polymorphic language. *Proceedings of the ACM on Programming Languages* 2, POPL (Jan. 2018), 1–29. <https://doi.org/10.1145/3158093> arXiv:1710.09756 [cs].
- Frédéric Bour, Basile Clément, and Gabriel Scherer. 2021. Tail Modulo Cons. In *JFLA 2021 - Journées Francophones des Langages Applicatifs*. Saint Médard d'Excideuil, France. <https://inria.hal.science/hal-03146495>
- Henry DeYoung, Frank Pfenning, and Klaas Pruiksma. 2020. Semi-Axiomatic Sequent Calculus. In *5th International Conference on Formal Structures for Computation and Deduction (FSCD 2020) (Leibniz International Proceedings in Informatics (LIPIcs), Vol. 167)*, Zena M. Ariola (Ed.). Schloss Dagstuhl – Leibniz-Zentrum für Informatik, Dagstuhl, Germany, 29:1–29:22. <https://doi.org/10.4230/LIPIcs.FSCD.2020.29>
- Jeremy Gibbons. 1993. Linear-time Breadth-first Tree Algorithms: An Exercise in the Arithmetic of Folds and Zips. No. 71 (1993). <https://www.cs.ox.ac.uk/publications/publication2363-abstract.html> Number: No. 71.
- Jeremy Gibbons, Donnacha Oisin Kidney, Tom Schrijvers, and Nicolas Wu. 2023. Phases in Software Architecture. In *Proceedings of the 1st ACM SIGPLAN International Workshop on Functional Software Architecture*. ACM, Seattle WA USA, 29–33. <https://doi.org/10.1145/3609025.3609479>
- J.-Y. Girard. 1995. Linear Logic: its syntax and semantics. In *Advances in Linear Logic*, Jean-Yves Girard, Yves Lafont, and Laurent Regnier (Eds.). Cambridge University Press, Cambridge, 1–42. <https://doi.org/10.1017/CBO9780511629150.002>
- Robert Hood and Robert Melville. 1981. Real-time queue operations in pure LISP. *Inform. Process. Lett.* 13, 2 (1981), 50–54. [https://doi.org/10.1016/0020-0190\(81\)90030-2](https://doi.org/10.1016/0020-0190(81)90030-2)
- John Hughes. 1986. A Novel Representation of Lists and its Application to the Function "reverse". *Inf. Process. Lett.* 22 (01 1986), 141–144.
- Yasuhiko Minamide. 1998. A functional representation of data structures with a hole. In *Proceedings of the 25th ACM SIGPLAN-SIGACT symposium on Principles of programming languages (POPL '98)*. Association for Computing Machinery, New York, NY, USA, 75–84. <https://doi.org/10.1145/268946.268953>
- Francesco Zappa Nardelli. 2009. Locally nameless backend for Ott. https://fzn.fr/projects/ln_ott/
- Chris Okasaki. 2000. Breadth-first numbering: lessons from a small exercise in algorithm design. In *Proceedings of the fifth ACM SIGPLAN international conference on Functional programming (ICFP '00)*. Association for Computing Machinery, New York, NY, USA, 131–136. <https://doi.org/10.1145/351240.351253>
- Peter Sewell, Francesco Zappa Nardelli, Scott Owens, Gilles Peskine, Thomas Ridge, Susmit Sarkar, and Rok Strniša. 2007. Ott: effective tool support for the working semanticist. In *Proceedings of the 12th ACM SIGPLAN International Conference on Functional Programming (Freiburg, Germany) (ICFP '07)*. Association for Computing Machinery, New York, NY, USA, 1–12. <https://doi.org/10.1145/1291151.1291155>
- Amir Shaikhha, Andrew Fitzgibbon, Simon Peyton Jones, and Dimitrios Vytiniotis. 2017. Destination-passing style for efficient memory management. In *Proceedings of the 6th ACM SIGPLAN International Workshop on Functional High-Performance Computing*. ACM, Oxford UK, 12–23. <https://doi.org/10.1145/3122948.3122949>
- Arnaud Spiwack, Csongor Kiss, Jean-Philippe Bernardy, Nicolas Wu, and Richard A. Eisenberg. 2022. Linearly qualified types: generic inference for capabilities and uniqueness. *Proceedings of the ACM on Programming Languages* 6, ICFP (Aug. 2022), 95:137–95:164. <https://doi.org/10.1145/3547626>

A FULL REDUCTION RULES FOR λ_d

$$\begin{array}{c}
\frac{\text{NotVal } t}{C[t' \ t] \rightarrow (C \circ (t' \ [])) [t]} \text{Focus-App1} \qquad \frac{}{(C \circ (t' \ [])) [v] \rightarrow C[t' \ v]} \text{UNFOCUS-App1} \\
\frac{\text{NotVal } t'}{C[t' \ v] \rightarrow (C \circ ([\ v])) [t']} \text{Focus-App2} \qquad \frac{}{(C \circ ([\ v])) [v'] \rightarrow C[v' \ v]} \text{UNFOCUS-App2} \\
\frac{}{C[(\lambda x_{\text{m}} \mapsto u) \ v] \rightarrow C[u[x := v]]} \text{RED-APP} \qquad \frac{\text{NotVal } t}{C[t \ ; \ u] \rightarrow (C \circ ([\ ; \ u])) [t]} \text{Focus-PatU} \\
\frac{}{(C \circ ([\ ; \ u])) [v] \rightarrow C[v \ ; \ u]} \text{UNFOCUS-PatU} \qquad \frac{}{C[(\ ; \ u] \rightarrow C[u]} \text{RED-PatU} \\
\frac{\text{NotVal } t}{C[\text{case}_{\text{m}} \ t \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}] \rightarrow (C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}))[t]} \text{Focus-PATS} \\
\frac{}{(C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}))[v] \rightarrow C[\text{case}_{\text{m}} \ v \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}]} \text{UNFOCUS-PATS} \\
\frac{}{C[\text{case}_{\text{m}} \ (\text{Inl } v_1) \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}] \rightarrow C[u_1[x_1 := v_1]]} \text{RED-PATL} \\
\frac{}{C[\text{case}_{\text{m}} \ (\text{Inr } v_2) \ \text{of} \ \{\text{Inl } x_1 \mapsto u_1, \text{Inr } x_2 \mapsto u_2\}] \rightarrow C[u_2[x_2 := v_2]]} \text{RED-PATR} \\
\frac{\text{NotVal } t}{C[\text{case}_{\text{m}} \ t \ \text{of} \ (x_1, x_2) \mapsto u] \rightarrow (C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ (x_1, x_2) \mapsto u)) [t]} \text{Focus-PATP} \\
\frac{}{(C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ (x_1, x_2) \mapsto u)) [v] \rightarrow C[\text{case}_{\text{m}} \ v \ \text{of} \ (x_1, x_2) \mapsto u]} \text{UNFOCUS-PATP} \\
\frac{}{C[\text{case}_{\text{m}} \ (v_1, v_2) \ \text{of} \ (x_1, x_2) \mapsto u] \rightarrow C[u[x_1 := v_1][x_2 := v_2]]} \text{RED-PATP} \\
\frac{\text{NotVal } t}{C[\text{case}_{\text{m}} \ t \ \text{of} \ E_{\text{n}} \ x \mapsto u] \rightarrow (C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ E_{\text{n}} \ x \mapsto u)) [t]} \text{Focus-PATE} \\
\frac{}{(C \circ (\text{case}_{\text{m}} \ [] \ \text{of} \ E_{\text{n}} \ x \mapsto u)) [v] \rightarrow C[\text{case}_{\text{m}} \ v \ \text{of} \ E_{\text{n}} \ x \mapsto u]} \text{UNFOCUS-PATE} \\
\frac{}{C[\text{case}_{\text{m}} \ E_{\text{n}} \ v' \ \text{of} \ E_{\text{n}} \ x \mapsto u] \rightarrow C[u[x := v']]} \text{RED-PATE} \\
\frac{\text{NotVal } t}{C[\text{map } t \ \text{with } x \mapsto t'] \rightarrow (C \circ (\text{map } [] \ \text{of } x \mapsto t')) [t]} \text{Focus-MAP} \\
\frac{}{(C \circ (\text{map } [] \ \text{of } x \mapsto t')) [v] \rightarrow C[\text{map } v \ \text{with } x \mapsto t']} \text{UNFOCUS-MAP} \\
\frac{h'' = \max(\text{hnames}(C)) + 1}{C[\text{map }_{\text{H}} \langle v_2 \ \wedge \ v_1 \rangle \ \text{with } x \mapsto t'] \rightarrow (C \circ ({}^{\text{OP}}_{H \pm h''} \langle v_2 [H \pm h''] \ \wedge \ [] \rangle)) [t' [x := v_1 [H \pm h'']]]} \text{OPEN-AMPAR} \\
\frac{}{(C \circ {}^{\text{OP}}_H \langle v_2 \ \wedge \ [] \rangle) [v_1] \rightarrow C[H \langle v_2 \ \wedge \ v_1 \rangle]} \text{CLOSE-AMPAR}
\end{array}$$

Fig. 12. Full reduction rules for λ_d (part 1)

$$\begin{array}{c}
\text{NotVal } u \\
\frac{}{C[\mathbf{to}_K u] \rightarrow (C \circ (\mathbf{to}_K [])) [u]} \text{Focus-ToA} \qquad \frac{}{(C \circ (\mathbf{to}_K [])) [v_2] \rightarrow C[\mathbf{to}_K v_2]} \text{Unfocus-ToA} \\
\\
\frac{}{C[\mathbf{to}_K v_2] \rightarrow C[\{\} \langle v_2 \wedge \rangle]} \text{RED-ToA} \qquad \frac{\text{NotVal } t}{C[\mathbf{from}_K t] \rightarrow (C \circ (\mathbf{from}_K [])) [t]} \text{Focus-FROMA} \\
\\
\frac{}{(C \circ (\mathbf{from}_K [])) [v] \rightarrow C[\mathbf{from}_K v]} \text{Unfocus-FROMA} \\
\\
\frac{}{C[\mathbf{from}_K \{\} \langle v_2 \wedge E_{\text{too}} v_1 \rangle] \rightarrow C[(v_2, E_{\text{too}} v_1)]} \text{RED-FROMA} \qquad \frac{\text{NotVal } t}{C[t \triangleleft ()] \rightarrow (C \circ ([\triangleleft ()])) [t]} \text{Focus-FILLU} \\
\\
\frac{}{(C \circ ([\triangleleft ()])) [v] \rightarrow C[v \triangleleft ()]} \text{Unfocus-FILLU} \qquad \frac{}{C[\rightarrow h \triangleleft ()] \rightarrow C[h := \{\} ()] [()]} \text{RED-FILLU} \\
\\
\frac{\text{NotVal } t}{C[t \triangleleft \text{Inl}] \rightarrow (C \circ ([\triangleleft \text{Inl}])) [t]} \text{Focus-FILLL} \qquad \frac{}{(C \circ ([\triangleleft \text{Inl}])) [v] \rightarrow C[v \triangleleft \text{Inl}]} \text{Unfocus-FILLL} \\
\\
\frac{h' = \max(\text{hnames}(C) \cup \{h\}) + 1}{C[\rightarrow h \triangleleft \text{Inl}] \rightarrow C[h := \{h'+1\} \text{Inl} [\overline{h'+1}] [\rightarrow h'+1]]} \text{RED-FILLL} \qquad \frac{\text{NotVal } t}{C[t \triangleleft \text{Inr}] \rightarrow (C \circ ([\triangleleft \text{Inr}])) [t]} \text{Focus-FILLR} \\
\\
\frac{}{(C \circ ([\triangleleft \text{Inr}])) [v] \rightarrow C[v \triangleleft \text{Inr}]} \text{Unfocus-FILLR} \\
\\
\frac{h' = \max(\text{hnames}(C) \cup \{h\}) + 1}{C[\rightarrow h \triangleleft \text{Inr}] \rightarrow C[h := \{h'+1\} \text{Inr} [\overline{h'+1}] [\rightarrow h'+1]]} \text{RED-FILLR} \qquad \frac{\text{NotVal } t}{C[t \triangleleft E_m] \rightarrow (C \circ ([\triangleleft E_m])) [t]} \text{Focus-FILLE} \\
\\
\frac{}{(C \circ ([\triangleleft E_m])) [v] \rightarrow C[v \triangleleft E_m]} \text{Unfocus-FILLE} \\
\\
\frac{h' = \max(\text{hnames}(C) \cup \{h\}) + 1}{C[\rightarrow h \triangleleft E_m] \rightarrow C[h := \{h'+1\} E_m [\overline{h'+1}] [\rightarrow h'+1]]} \text{RED-FILLE} \qquad \frac{\text{NotVal } t}{C[t \triangleleft (,)] \rightarrow (C \circ ([\triangleleft (,)])) [t]} \text{Focus-FILLP} \\
\\
\frac{}{(C \circ ([\triangleleft (,)])) [v] \rightarrow C[v \triangleleft (,)]} \text{Unfocus-FILLP} \\
\\
\frac{h' = \max(\text{hnames}(C) \cup \{h\}) + 1}{C[\rightarrow h \triangleleft (,)] \rightarrow C[h := \{h'+1, h'+2\} (\overline{h'+1}, \overline{h'+2})][(\rightarrow h'+1, \rightarrow h'+2)]} \text{RED-FILLP} \\
\\
\frac{\text{NotVal } t}{C[t \triangleleft (\lambda x_m \mapsto u)] \rightarrow (C \circ ([\triangleleft (\lambda x_m \mapsto u)])) [t]} \text{Focus-FILLF} \\
\\
\frac{}{(C \circ ([\triangleleft (\lambda x_m \mapsto u)])) [v] \rightarrow C[v \triangleleft (\lambda x_m \mapsto u)]} \text{Unfocus-FILLF} \\
\\
\frac{}{C[\rightarrow h \triangleleft (\lambda x_m \mapsto u)] \rightarrow C[h := \{\} \lambda x_m \mapsto u] [()]} \text{RED-FILLF} \\
\\
\frac{\text{NotVal } t}{C[t \triangleleft t'] \rightarrow (C \circ ([\triangleleft t'])) [t]} \text{Focus-FILLComp1} \qquad \frac{}{(C \circ ([\triangleleft t'])) [v] \rightarrow C[v \triangleleft t']} \text{Unfocus-FILLComp1} \\
\\
\frac{\text{NotVal } t'}{C[v \triangleleft t'] \rightarrow (C \circ (v \triangleleft [])) [t']} \text{Focus-FILLComp2} \qquad \frac{}{(C \circ (v \triangleleft [])) [v'] \rightarrow C[v \triangleleft v']} \text{Unfocus-FILLComp2} \\
\\
\frac{h' = \max(\text{hnames}(C) \cup \{h\}) + 1}{C[\rightarrow h \triangleleft H \langle v_2 \wedge v_1 \rangle] \rightarrow C[h := (H \triangleleft h') v_2 [H \triangleleft h']][v_1 [H \triangleleft h']]} \text{RED-FILLComp}
\end{array}$$

Fig. 13. Full reduction rules for λ_d (part 2)